

# Database Management Systems

- Optimization

*Doug Shook*

# How can we make things faster?

- Indexing
- Query Optimization

# Indexing

- Provides a secondary access path
  - Does not alter primary physical representation
- Indexes can take many forms:
  - Single Level
    - Primary, clustered
  - Multi Level
    - Trees

# Single Level

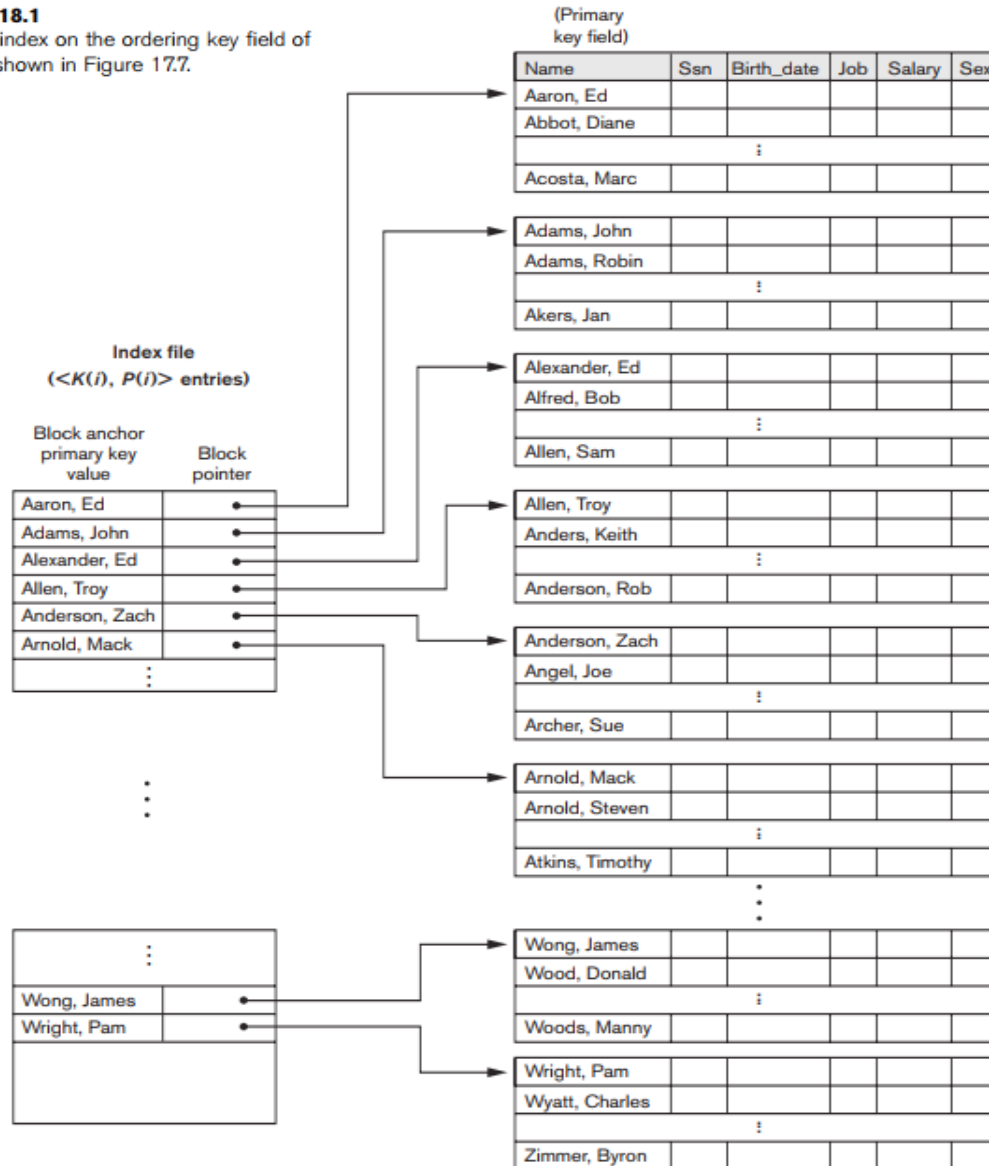
- Much like an index from a book
- Select a column to be indexed
  - Make a list of all values contained within that column
  - Create an index that has a column value as the key and a list of pages as the value
- Indexing values are ordered
  - How does this help?
- Size in comparison to primary storage?

# Primary Indexes

- One entry per page
- Key: First index value on the page
- Value: file address of the page
- Requires that data be stored in order (why?)
  - Requires that the index is based on the primary key (why?)
- Primary indexes are an example of sparse indexes
- Search time?

# Primary Indexes

**Figure 18.1**  
Primary index on the ordering key field of  
the file shown in Figure 17.7.



# Primary Indexes

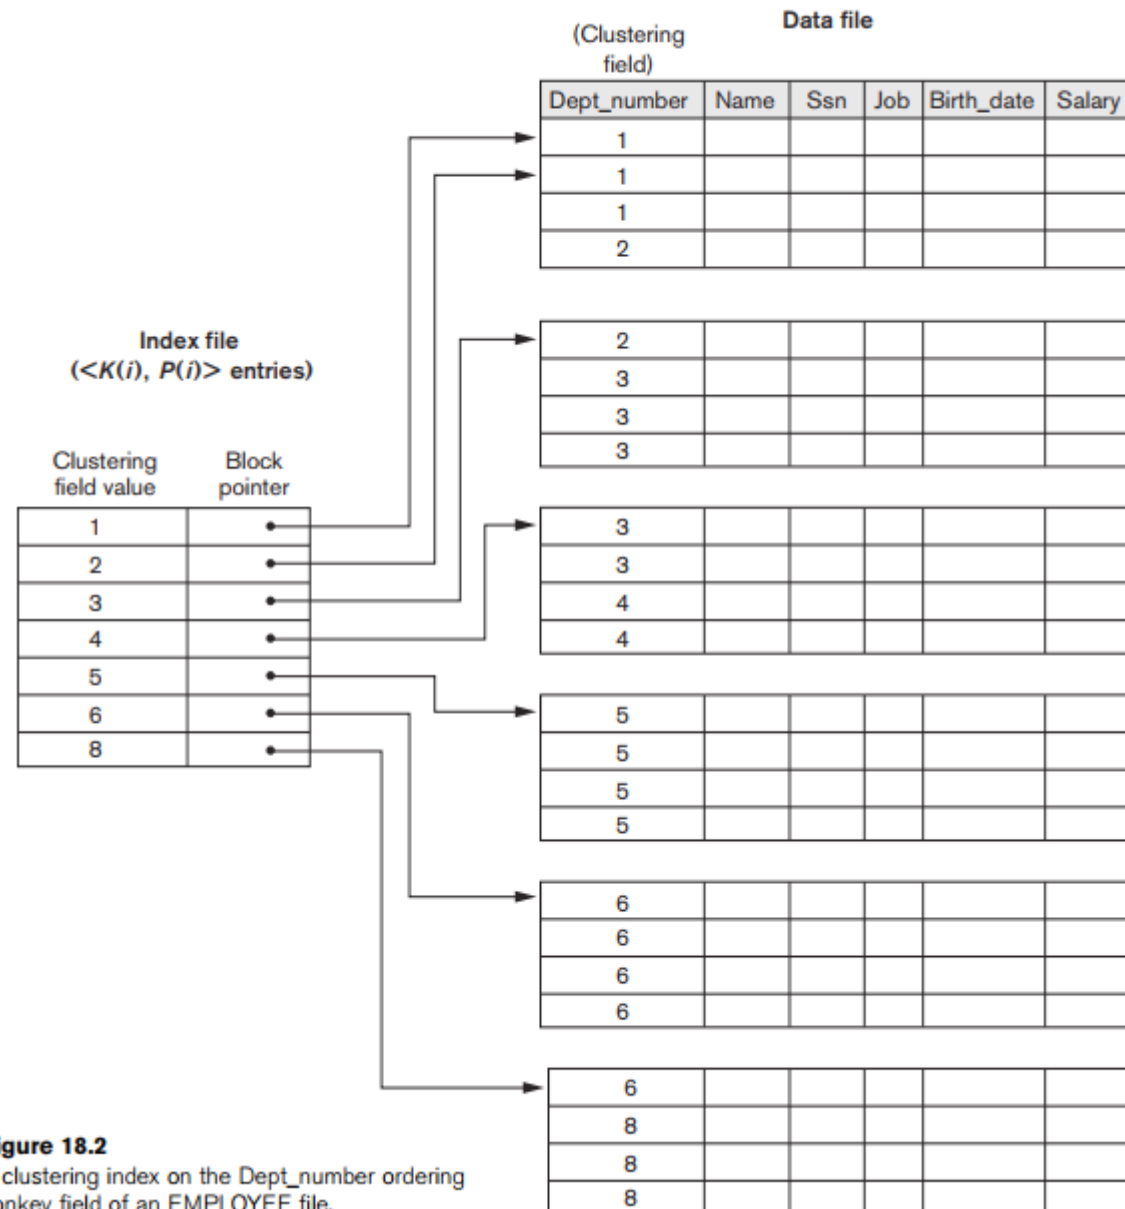
- How is this scheme affected by insertions and deletions?

# Clustered Index

- Indexing technique for non-key columns
  - What's the major difference?
- Is this sparse or dense
- How does this affect insertion and deletion?



# Clustered Index



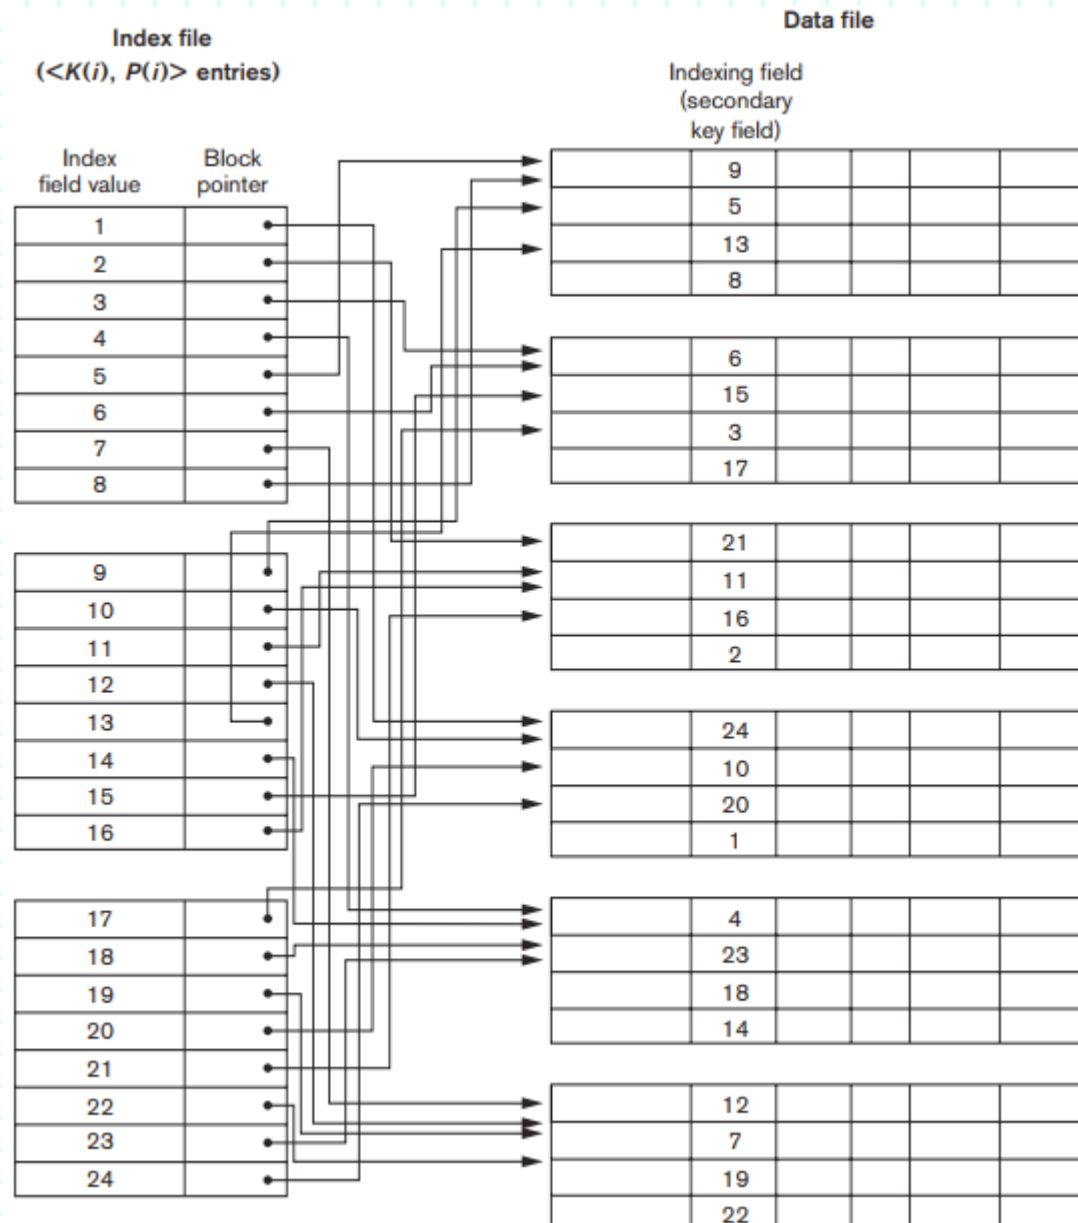
**Figure 18.2**

A clustering index on the Dept\_number ordering nonkey field of an EMPLOYEE file.

# Secondary Index

- What if we want to index by a column that is not ordered?
  - Unique vs. not unique

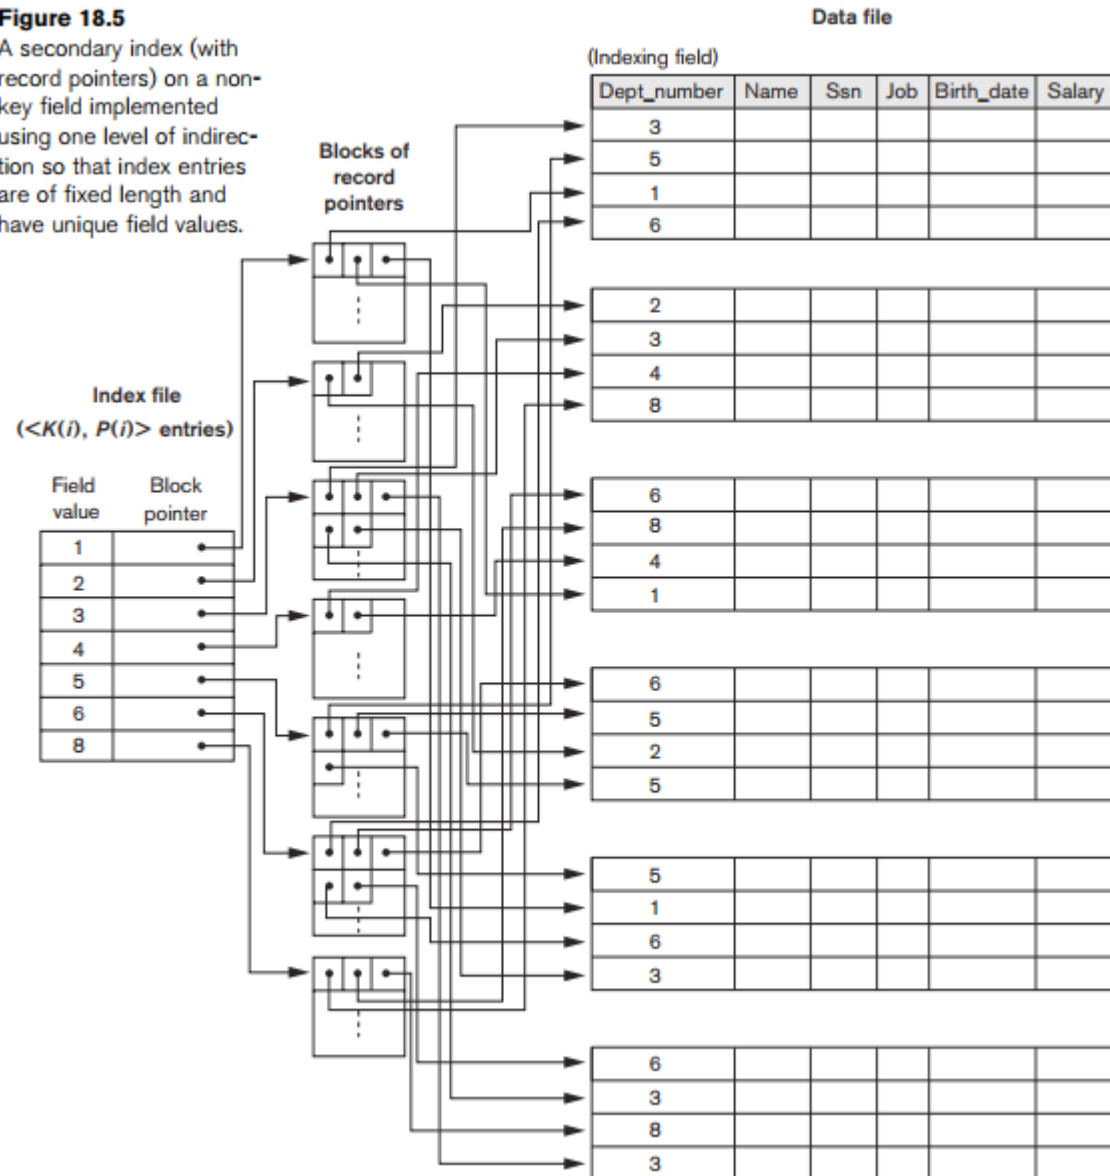
# Secondary Index - Unique



# Secondary – Not Unique

**Figure 18.5**

A secondary index (with record pointers) on a non-key field implemented using one level of indirection so that index entries are of fixed length and have unique field values.



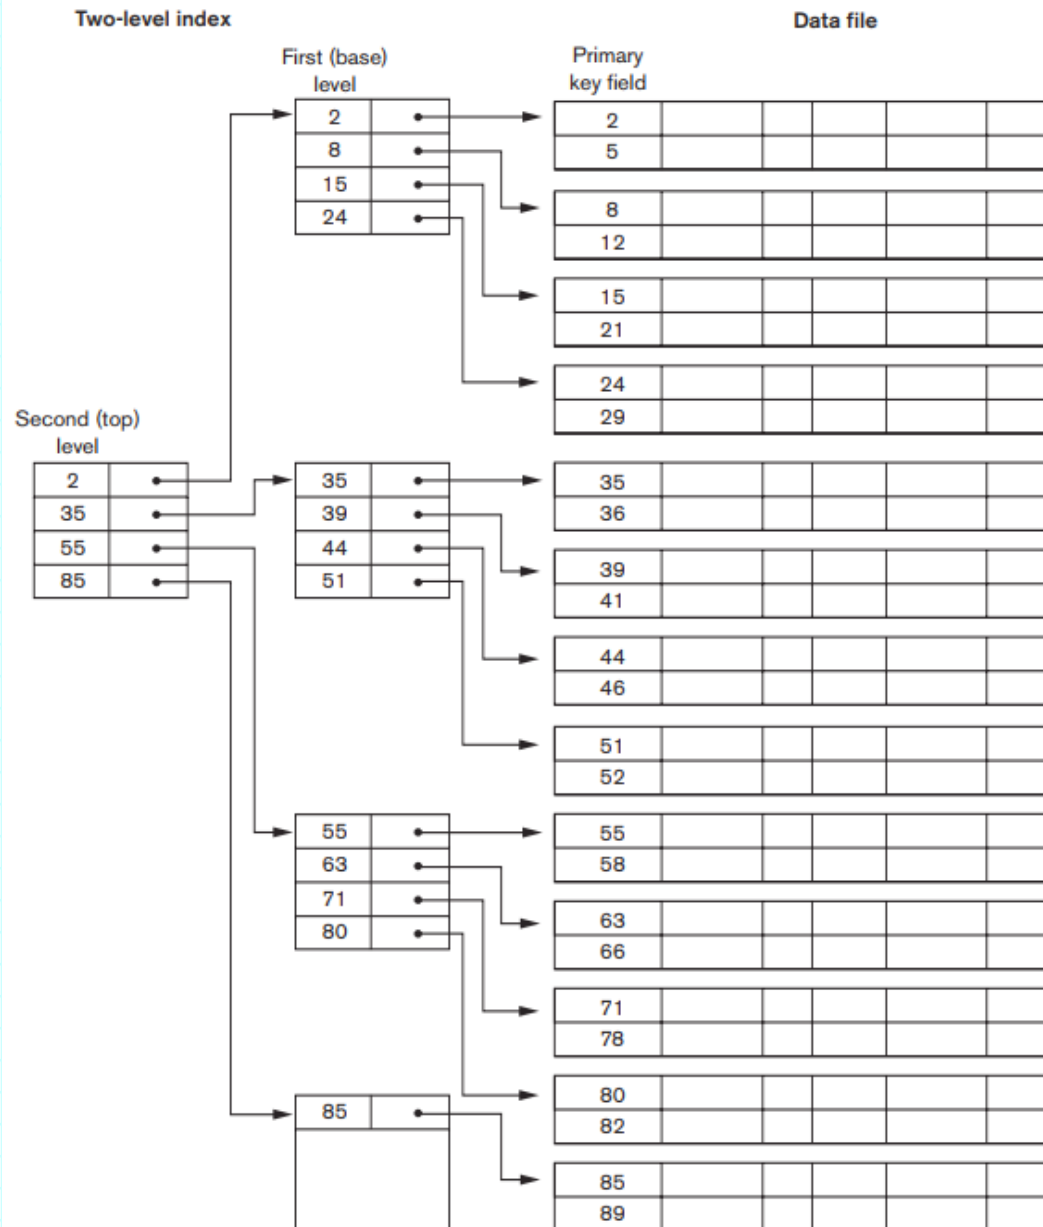
# Secondary Index

- Dense or sparse?
- Performance?
  - Performance improvement?

# Multilevel Indexes

- Same idea as a single level index
  - Try to reduce the search space even faster
- Idea: create an index (first layer)
  - Then create another index into that index
  - Repeat
- If we keep our indexes ordered, what kind of index can we use for the upper layers?
- Restrictions on index type of first layer?
- Limitations

# Multilevel Index



# Multilevel Index

- Search time?
- How to deal with insert and delete?



# Dynamic Multilevel Index

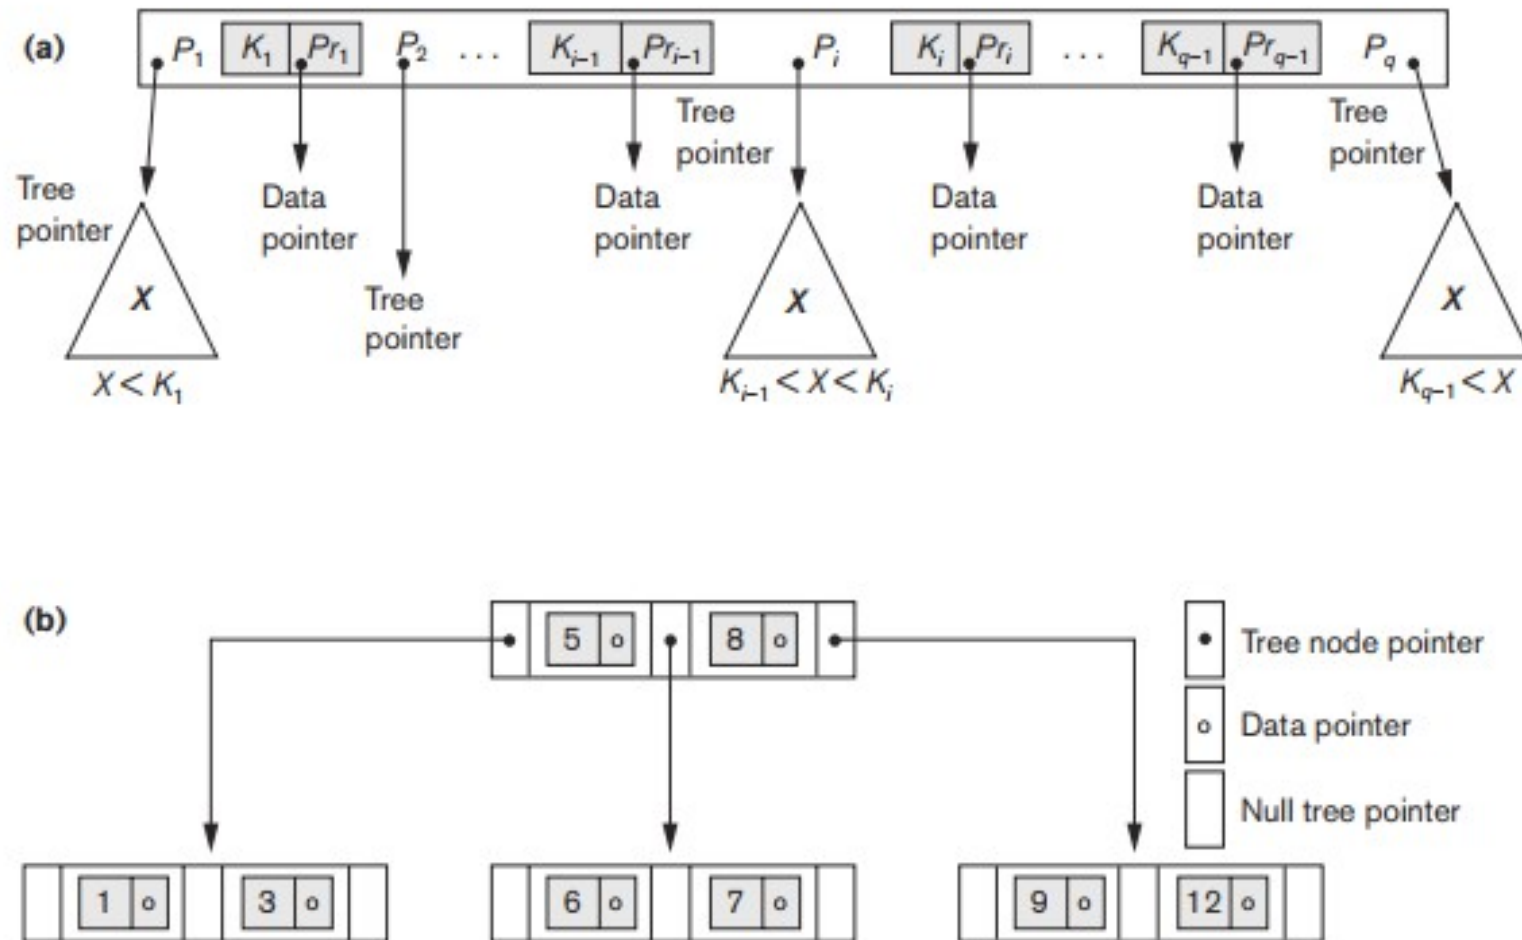
- Attempts to mitigate problems with insert and delete by leaving some empty space in each page of the index
  - Tradeoff?
- Uses search trees
  - B-Trees
  - B+-Trees

# B-Trees

- Properties
  - Always balanced
  - Tries to minimize wasted space due to deletions
  - Simplifies insertion and deletion (mostly)

# B-Trees

650 Chapter 18 Indexing Structures for Files



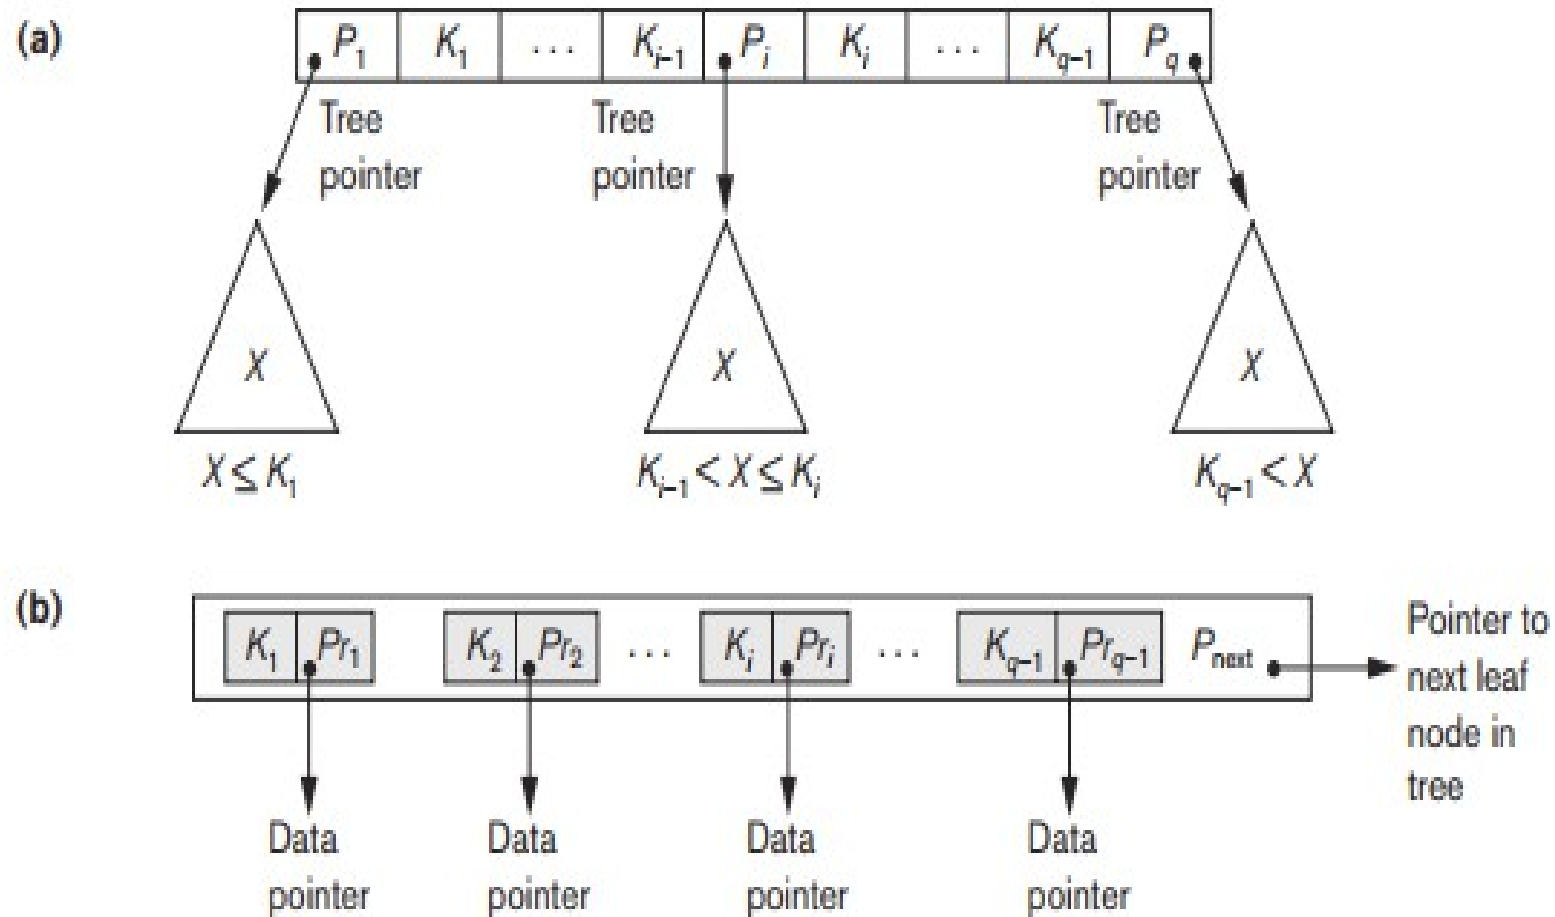
# B-Trees

- Previous example assumes we're searching a key
  - What if we're not?
- Insertion and deletion?

# B+-Trees

- Similar to a B-Tree
  - Data only stored at the leaf nodes

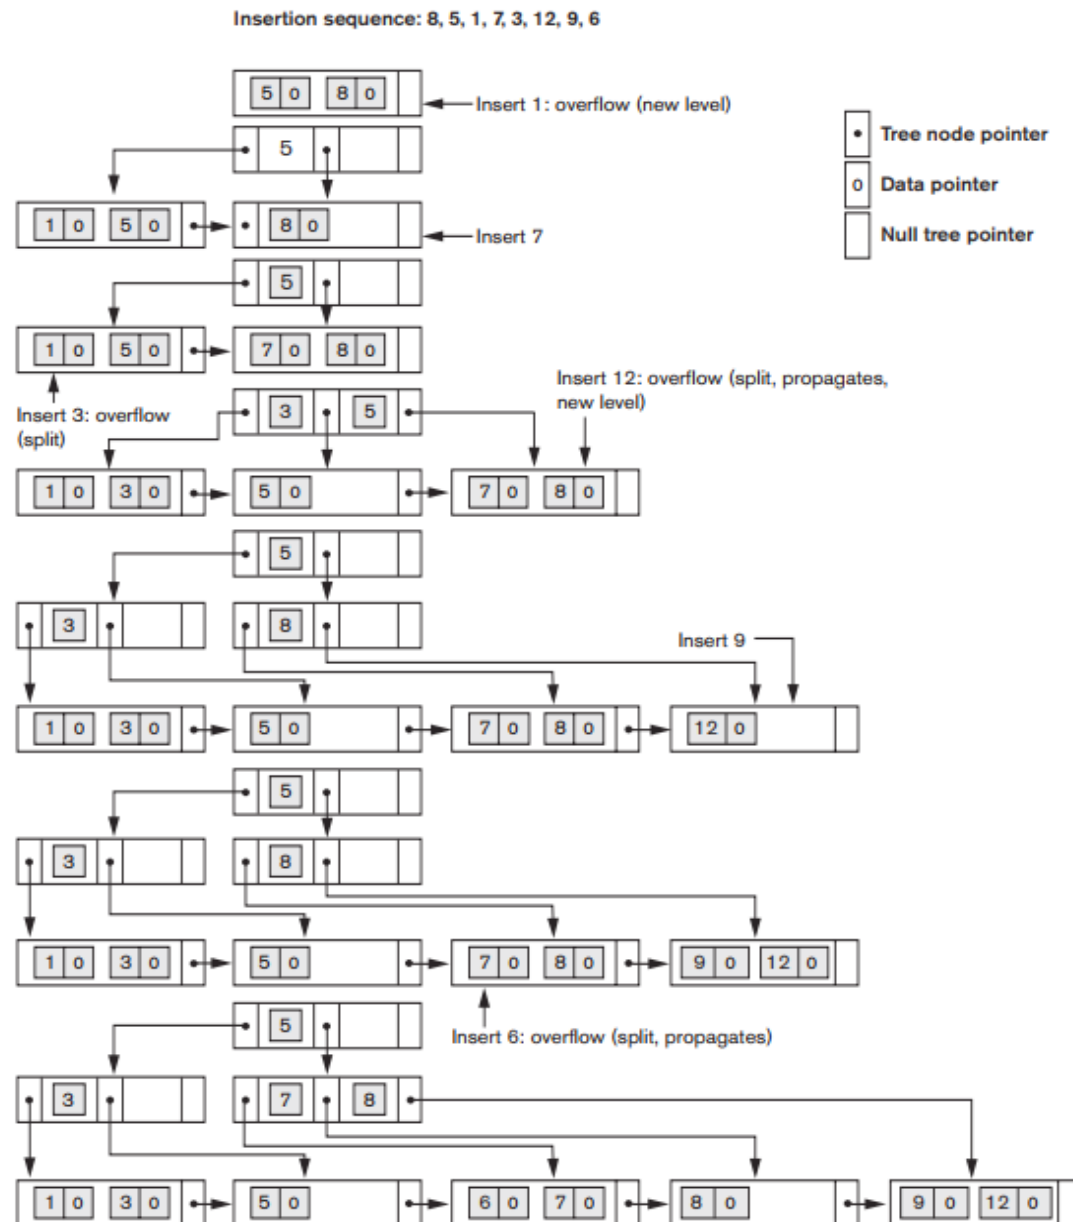
# B+-Trees



# B+-Trees

- Advantages of two different node types?
- Insertion and deletion?

# Insertion Example





# Deletion Example

