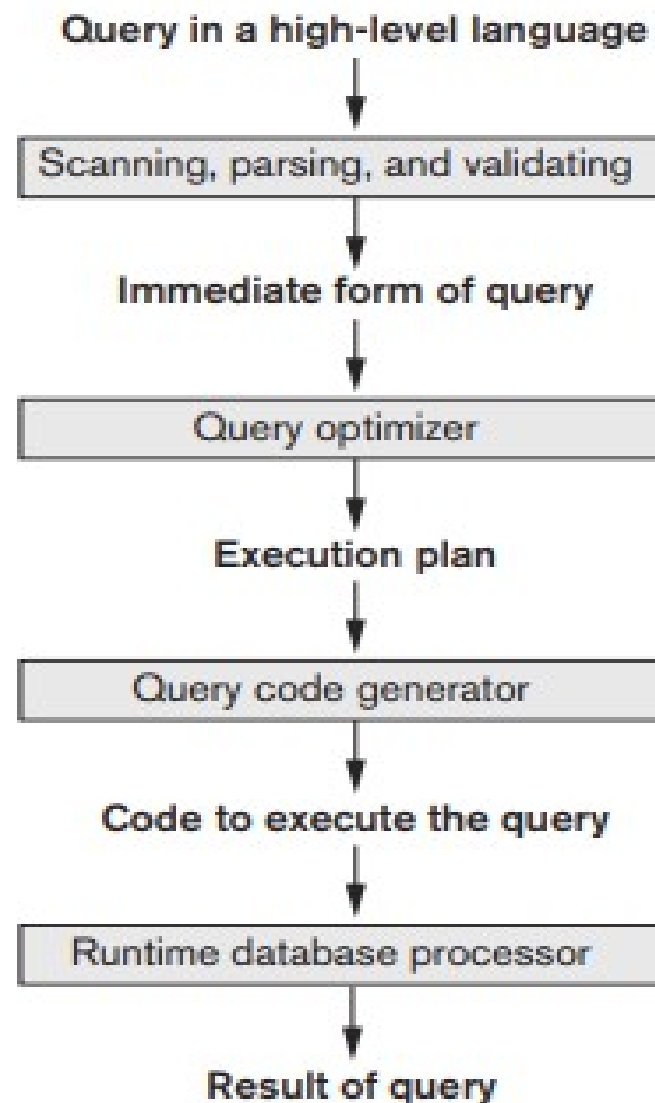


Database Management Systems

- Query Optimization

Doug Shook

Query Execution



Query Optimization

- Optimization is hard
 - Many possible combinations
- Maybe we aren't finding the “optimal” solution.....
- What operations should we focus on?

SELECT Optimization

- Consider a simple SELECT query:

```
SELECT  
FROM  
WHERE
```

- What possible approaches are there to executing this?
 - What if WHERE is more complex (uses AND)?

Selectivity

- Ratio of tuples that satisfy condition to the total number of records
- Can we compute this value exactly?
- How does this help us with optimization?

JOIN Optimization

- There's more than one way to perform a join
 - Nested Loop
 - Single Loop
 - Sort-merge
 - Partition-hash
- Which version is the simplest? Which version is the fastest?

Nested Loop Join

- How does table size affect this join?
- Example: assume one table has 10 pages, and another has 2000 pages. How many page accesses do we need for each nested loop configuration?

JOIN Selection Factor

- How many records from each table do we expect to match the join condition?
- How does this help us?
 - Which JOIN type does this affect?

Heuristics

- We can optimize queries by manipulating the query tree directly
 - Must satisfy order of operations
- Idea: one tree can be rewritten in numerous ways
 - Let's find the fastest version

Heuristics

- SELECT:
 - Can cascade conjunctive SELECT operations
 - Operations are commutative
- PROJECT
 - Cascading PROJECTS are somewhat irrelevant
 - Can be commuted with a SELECT
 - When?

Heuristics

■ JOINS

- Can be commuted with SELECT if select only affects one side of the join
- Can be commuted with PROJECT if all projected columns are part of the tables being joined

■ Set operations

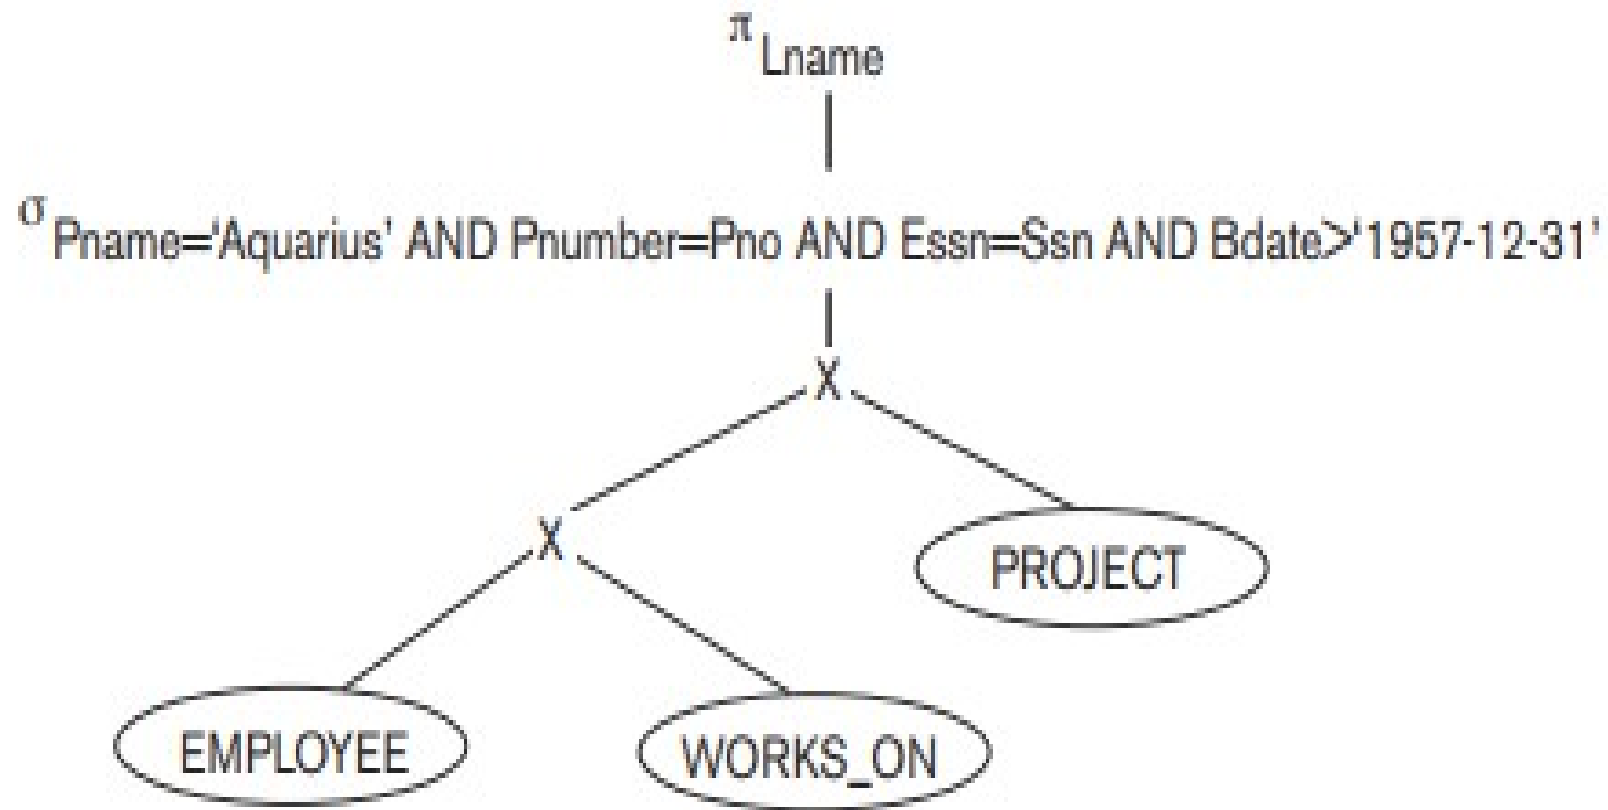
- Union and intersection are commutative
- Union, intersection, and join are associative
- All set operations are commutative with SELECT
- PROJECT is commutative with union

■ Cartesian product + SELECT = JOIN

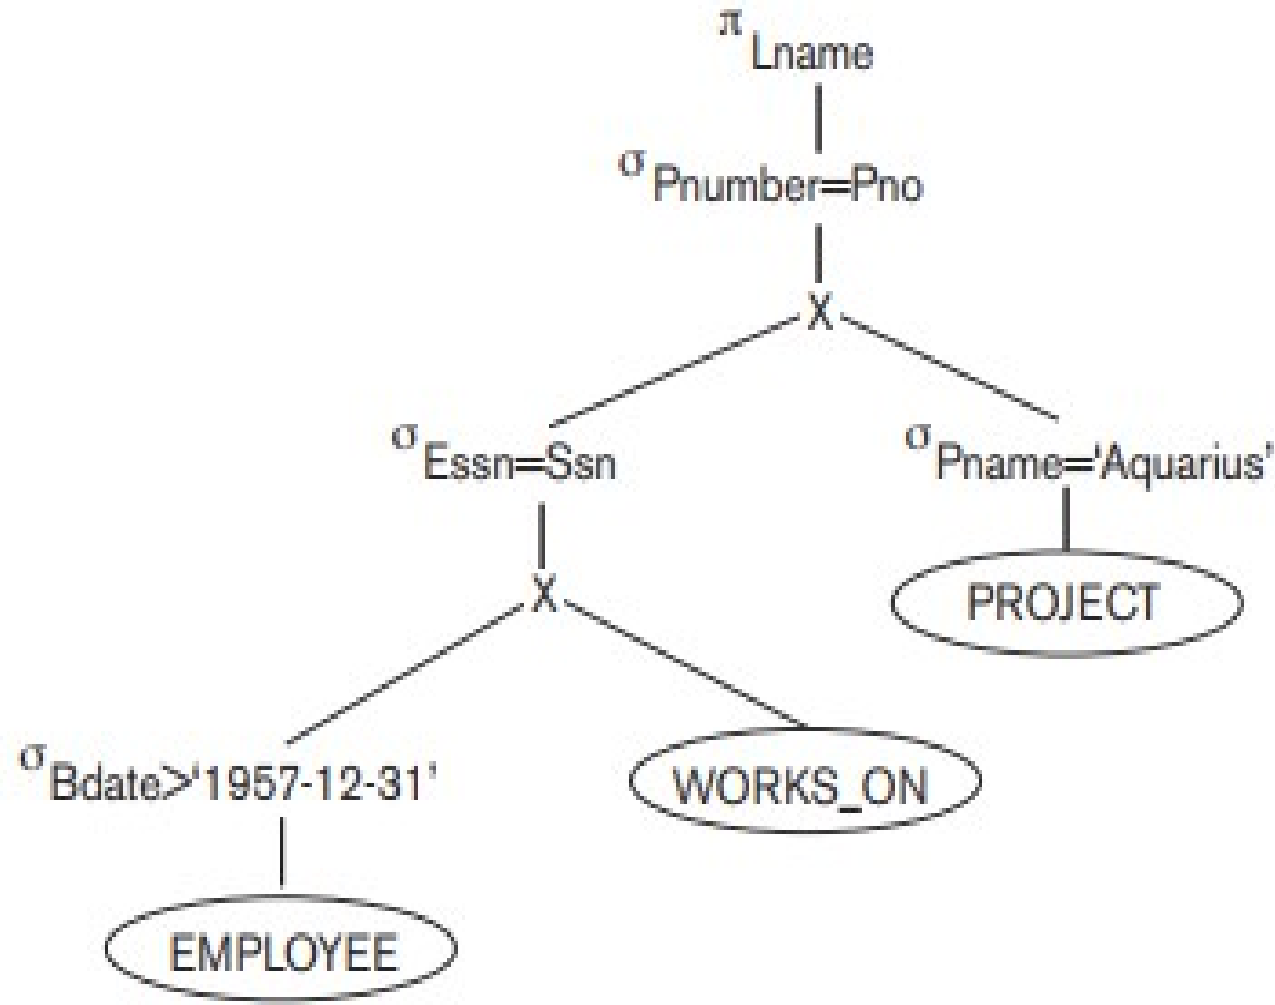
Optimization using Heuristics

- Break up conjunctive SELECTS
 - Allows us to move them around more easily
- Use SELECT commutativity to move SELECT operations as far down the tree as possible
- When performing operations on multiple tables, move relations with SELECT restrictions as far down the tree as possible
- Combine cartesian products with selects when applicable
- Move PROJECTs down the tree as far as possible

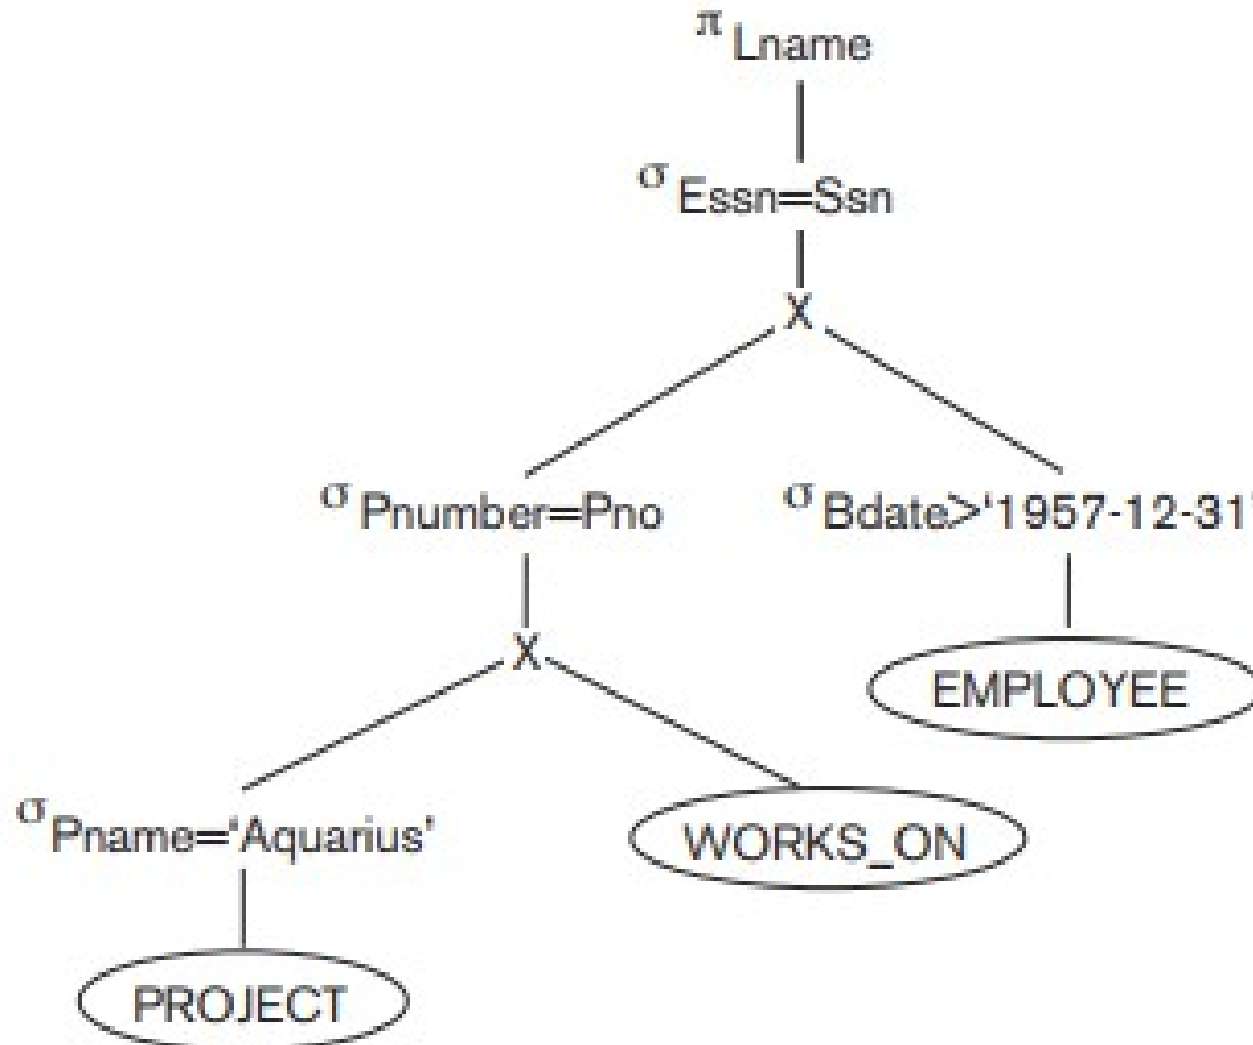
Heuristic Example



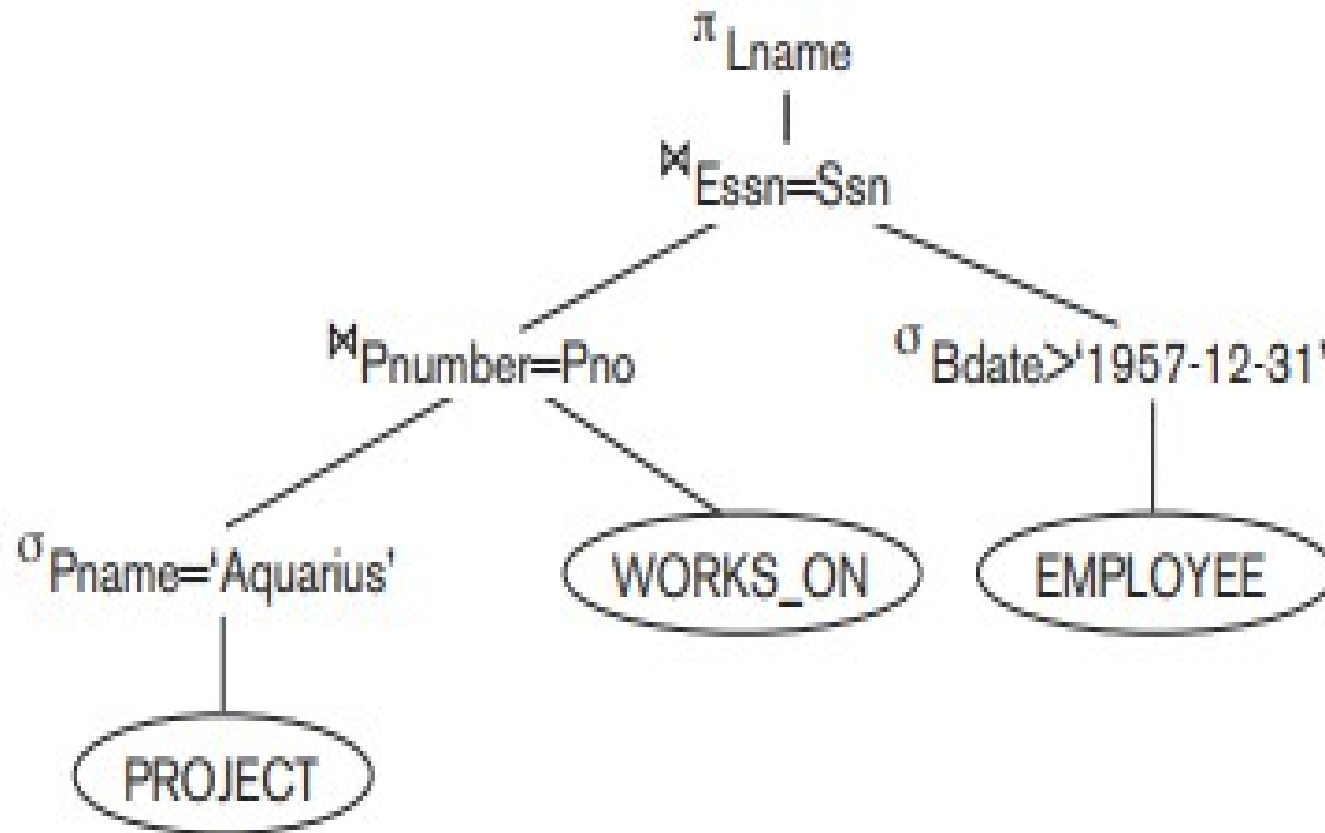
Heuristic Example



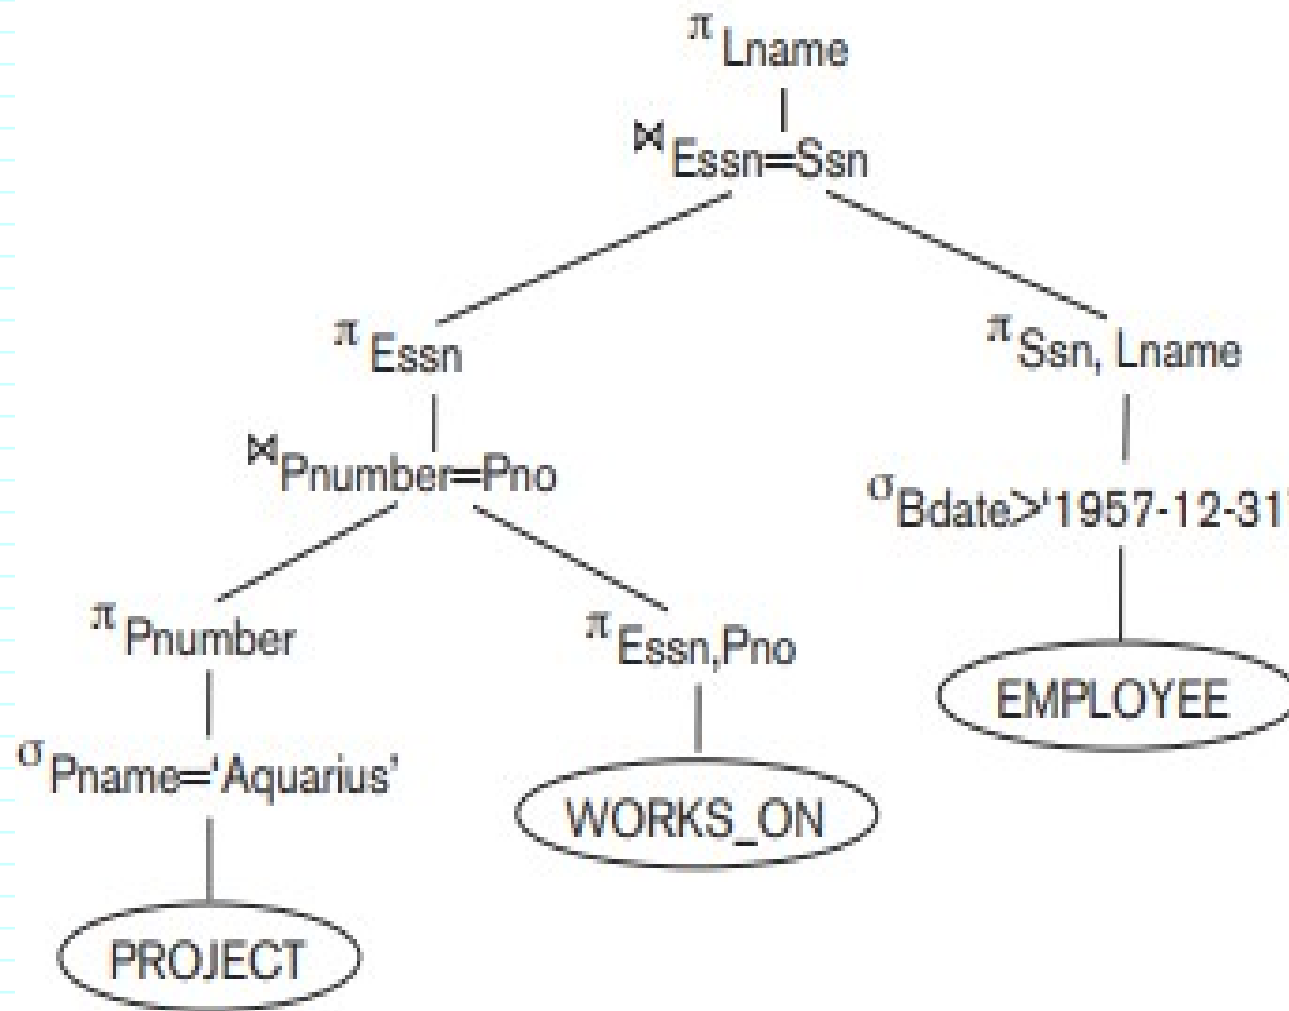
Heuristic Example



Heuristic Example



Heuristic Example



Heuristic Example

