

Database Management Systems

- Concurrency Control

Review

- Why do we need transactions?
 - What does a transaction contain?
 - What are the four properties of a transaction?
- What is a schedule?
 - What is a conflict?
 - What does it mean for a schedule to be serializable?

Locking

- We wish to create serializable schedules
- Many techniques for this:
 - Locking
 - Timestamps
 - Multiversion
- In practice, locking is used most frequently
 - Operates on data items
 - What is a data item?
 - Many types of locks available

Binary Locks

- Two states:
 - Locked/unlocked
 - Cannot access an item while it is locked
- Enforces mutual exclusion

Binary Locks

lock_item(X):

B: if $\text{LOCK}(X) = 0$ (* item is unlocked *)

then $\text{LOCK}(X) \leftarrow 1$ (* lock the item *)

else

begin

wait (until $\text{LOCK}(X) = 0$

and the lock manager wakes up the transaction);

go to **B**

end;

unlock_item(X):

$\text{LOCK}(X) \leftarrow 0;$ (* unlock the item *)

if any transactions are waiting

then wakeup one of the waiting transactions;

Implementation

- Must maintain a table of locks
 - And a queue for transactions that are waiting
- Must obey the following rules:
 - Must lock before any read or write
 - Must unlock after any read or write
 - Transaction cannot lock or unlock an item if it is already locked/unlocked
- In practice, while binary locks work just fine, they are considered too restrictive
 - Why?

Shared/Exclusive Locks

- We wish for our locks to mesh well with conflicts
 - Is it a problem for two transactions to read the same item?
- Two kinds of locks:
 - Read
 - Write
- Only one write lock at a time
 - Multiple read locks can exist for the same item
 - How to track them?

Shared/Exclusive Properties

- The following rules must be obeyed:
 - Must acquire a read or write lock before reading
 - Must acquire a write lock before writing
 - A transaction cannot reacquire a lock that it already has
- It is possible to convert locks
 - Read \rightarrow Write
 - Write \rightarrow Read
- What must be true for conversion to be allowed?

Two Phase Locking

- All locks should be acquired before the first unlock statement
 - Growing phase: increase number of locks
 - Shrinking phase: decrease number of locks
- Guarantees serializability

Two phase locking

T_1	T_2
<pre> read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>

T_1'	T_2'
<pre> read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>


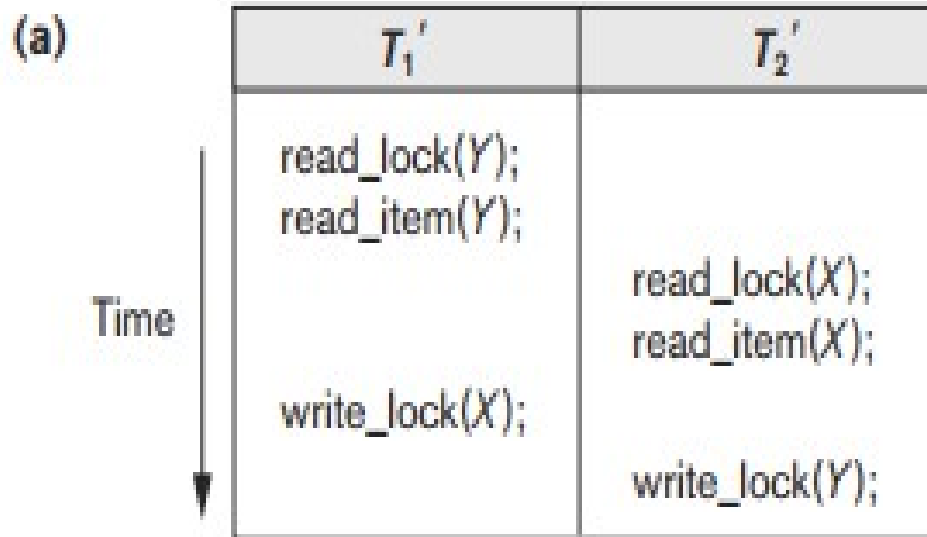
Types of Two Phase

- Basic
- Conservative
 - Must acquire all locks before beginning the transaction
- Strict
 - Do not release any write locks until after commit or abort
 - Leads to strict schedules
 - No transaction can read or write X until the last transaction that wrote X has committed

Issues With Two Phase

- Cannot generate all possible serializable schedules
- Reduces concurrency
- Can lead to deadlock
- Can lead to starvation

Deadlock



Deadlock Prevention

- Timestamp based
- Wait-die: older transaction is allowed to wait, younger transaction must abort and restart
- Wound-wait: Older transaction can abort younger transaction, restarting it later. Younger transactions simply wait.
- Can lead to unnecessary aborts

Deadlock Prevention

- Non-timestamp based
- No wait – if lock is unavailable, abort
- Cautious wait – If the transaction that has the lock is not waiting on another lock, then wait. Otherwise abort.

Deadlock Detection

- Create a graph
- Nodes are transactions
- Directed edges from transactions that wait on locks from other transactions
 - Drop edges when waiting is over
 - When do we know deadlock has occurred?
 - What do we do?
- Works best with small transactions
 - Why?

Starvation

- Occurs when a transaction has to wait a very long time.
 - Possibly indefinitely?
 - When does this happen?
- Solutions:
 - Time scaling priority
 - FCFS

Granularity

- What should we be locking?
 - Fields?
 - Columns?
 - Pages?
 - Tables?

- How does granularity affect concurrency?

- Why not just lock everything on the field level?

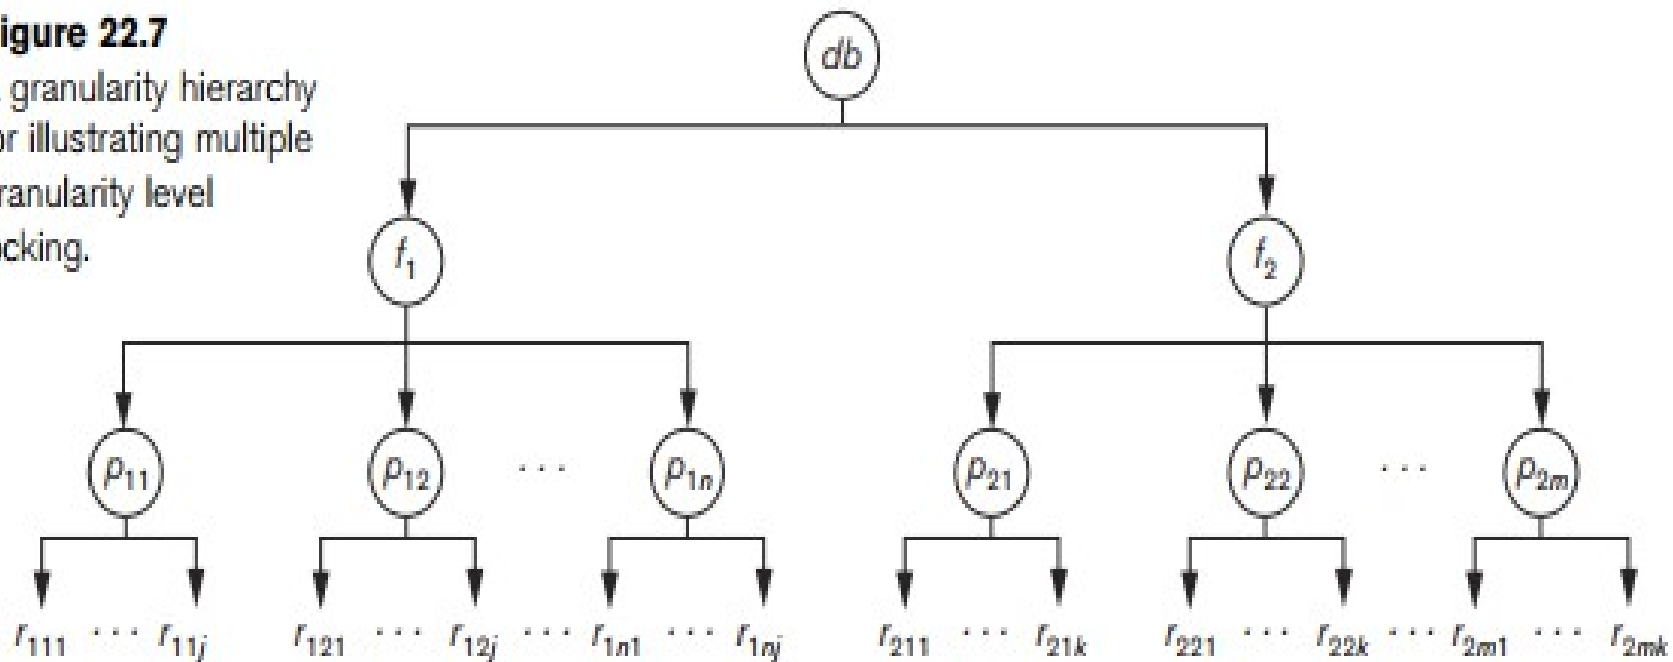
- What is the best size?

Granularity

- Two transactions: one wants to update all records in a file, the other wants to read a single record
 - How does order affect locking?

Figure 22.7

A granularity hierarchy for illustrating multiple granularity level locking.



Intention Locks

- What locks would be requested on descendants of the tree?
- Intention Shared
- Intention Exclusive
- Shared intention exclusive
 - Currently share-locked, intends to be exclusive

Intention Locks

- Lock the root first
- Can be locked in S or IS mode if parent is locked in IS or IX mode.
- Can be locked in X, IX, or SIX mode if parent is locked in IX or SIX mode
- Nodes can only be unlocked if children are unlocked

Intention Locks

	IS	IX	S	SIX	X
IS	Yes	Yes	Yes	Yes	No
IX	Yes	Yes	No	No	No
S	Yes	No	Yes	No	No
SIX	Yes	No	No	No	No
X	No	No	No	No	No

Intention Locks

- So for our previous example:
 - T1 locks root and file with IX
 - T2 may still lock records with IS/S while file is in IX
- If T2 goes first: root, file, record locked with IS/S
 - T1 cannot acquire lock until T2 is done
 - But the check is a lot faster!