

Database Management Systems

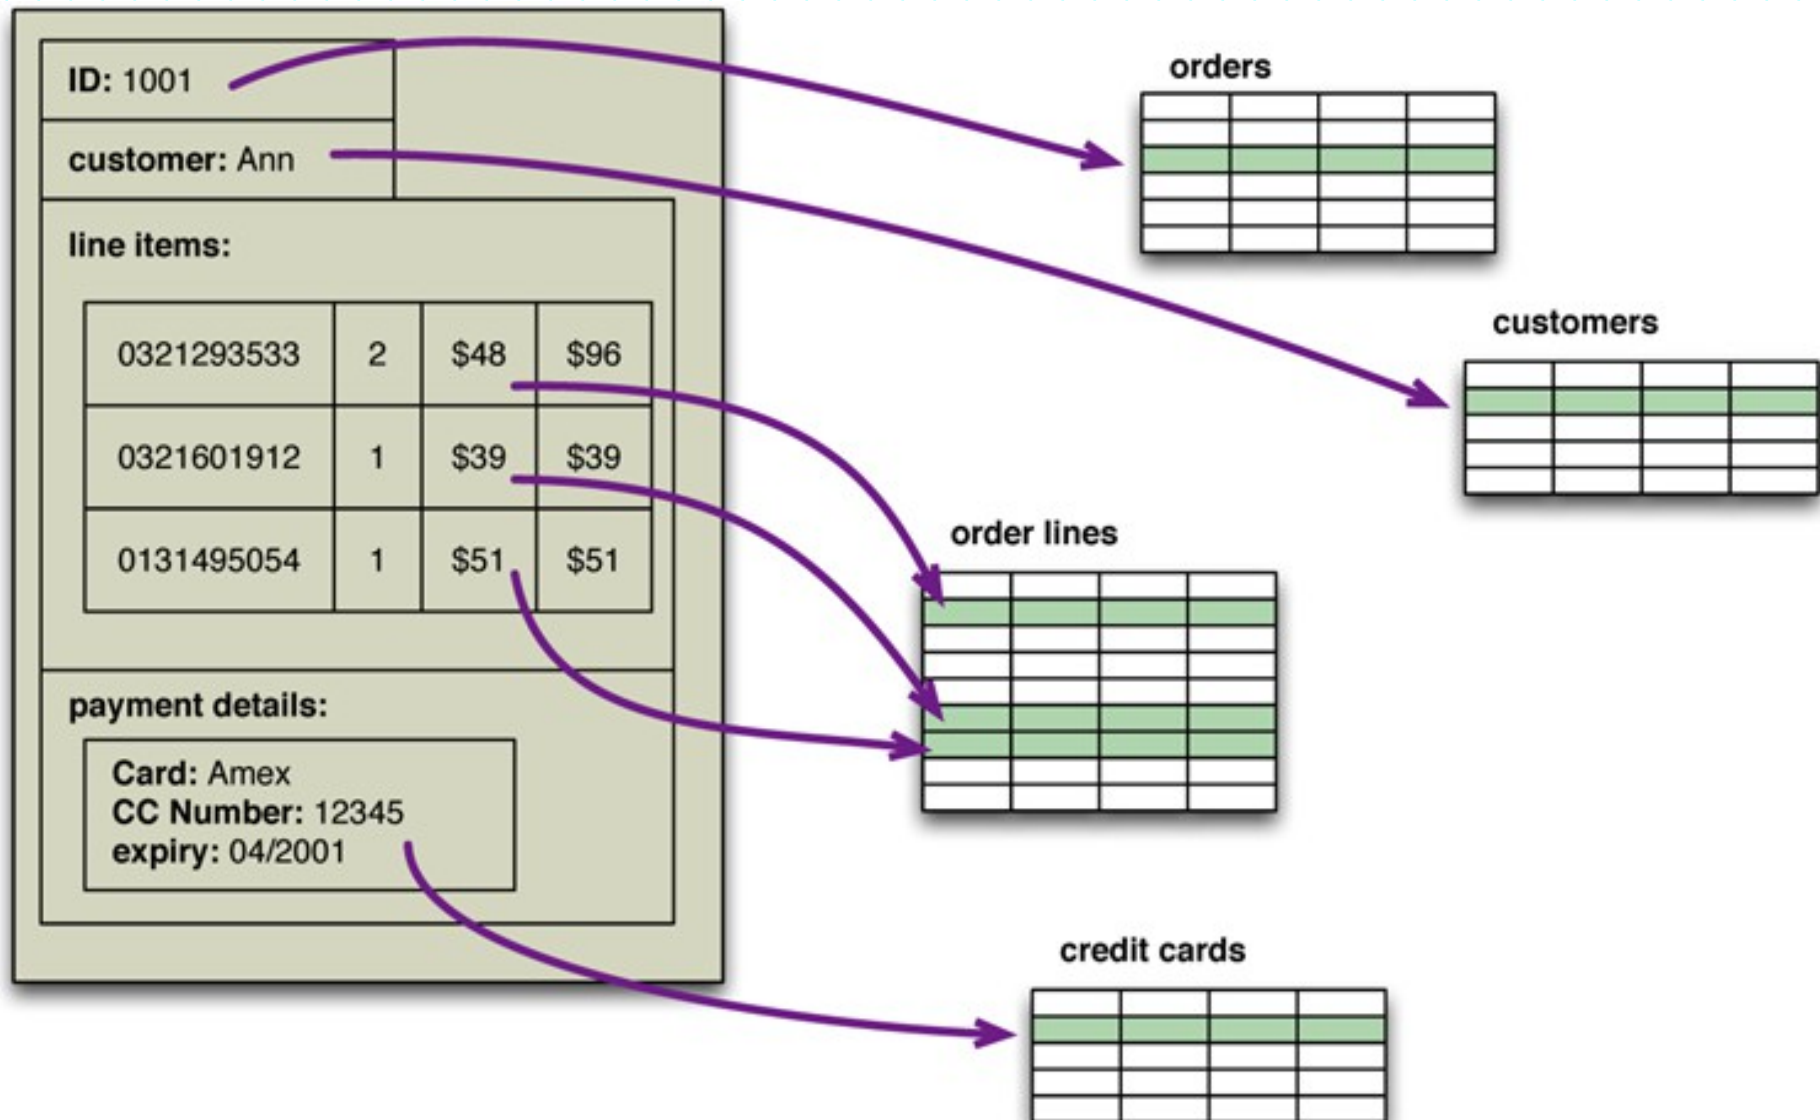
- NoSQL

Why NoSQL?

- Better question: why relational databases?
- Sadalage, Fowler, "NoSQL Distilled: A Brief Guide to the Emerging World of Polyglot Persistence" 2012

Problems with the Relational Model

■ Impedance Mismatch



Application vs. Integration

- What are you using the database for?
 - How does this effect...
 - Complexity?
 - Performance?
 - Development?

The effect of scale

- Consider how the internet looks now compared to 15 years ago....
 - How did we come to meet the increase in demand?

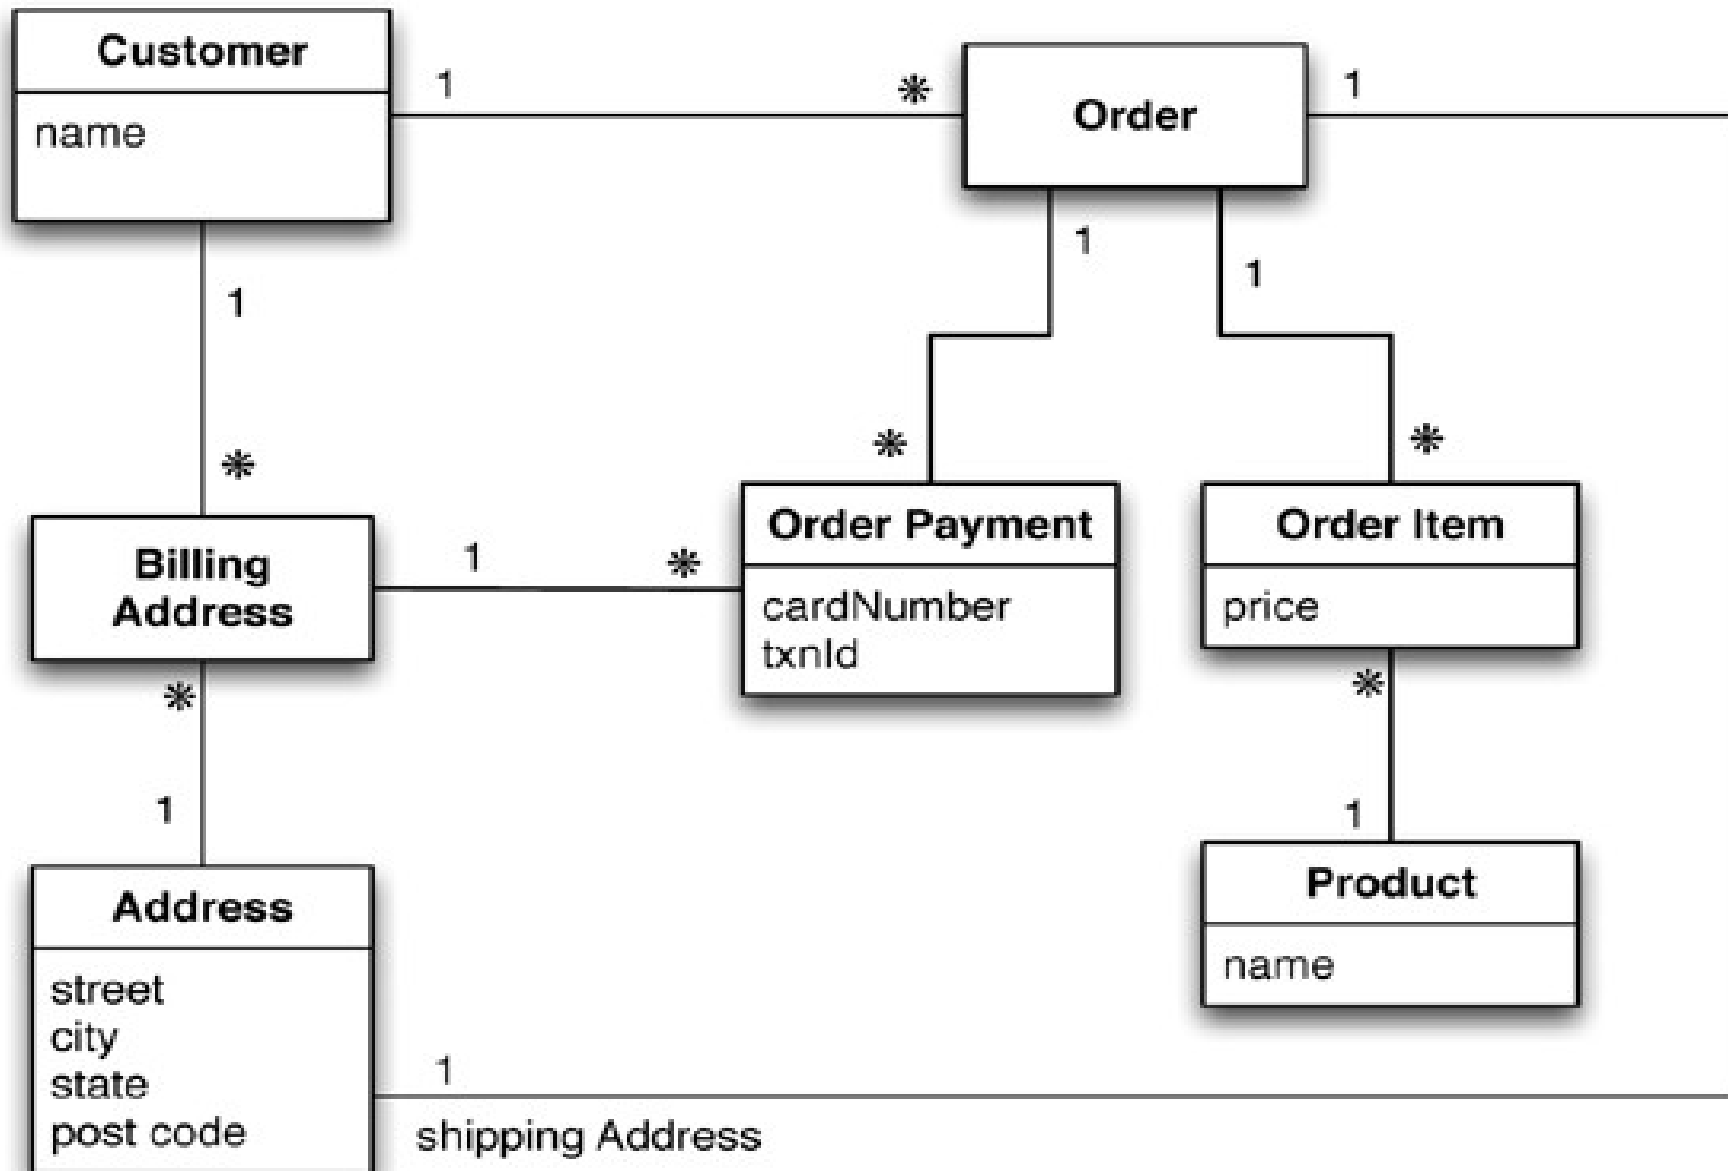
Here comes NoSQL

- What is it?
- How is it different?
- Why should we be using it?

Aggregates

- How does a relational database store records?
 - What are the limitations of this?

A relational data model



Relational data

Customer	
Id	Name
1	Martin

Order		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

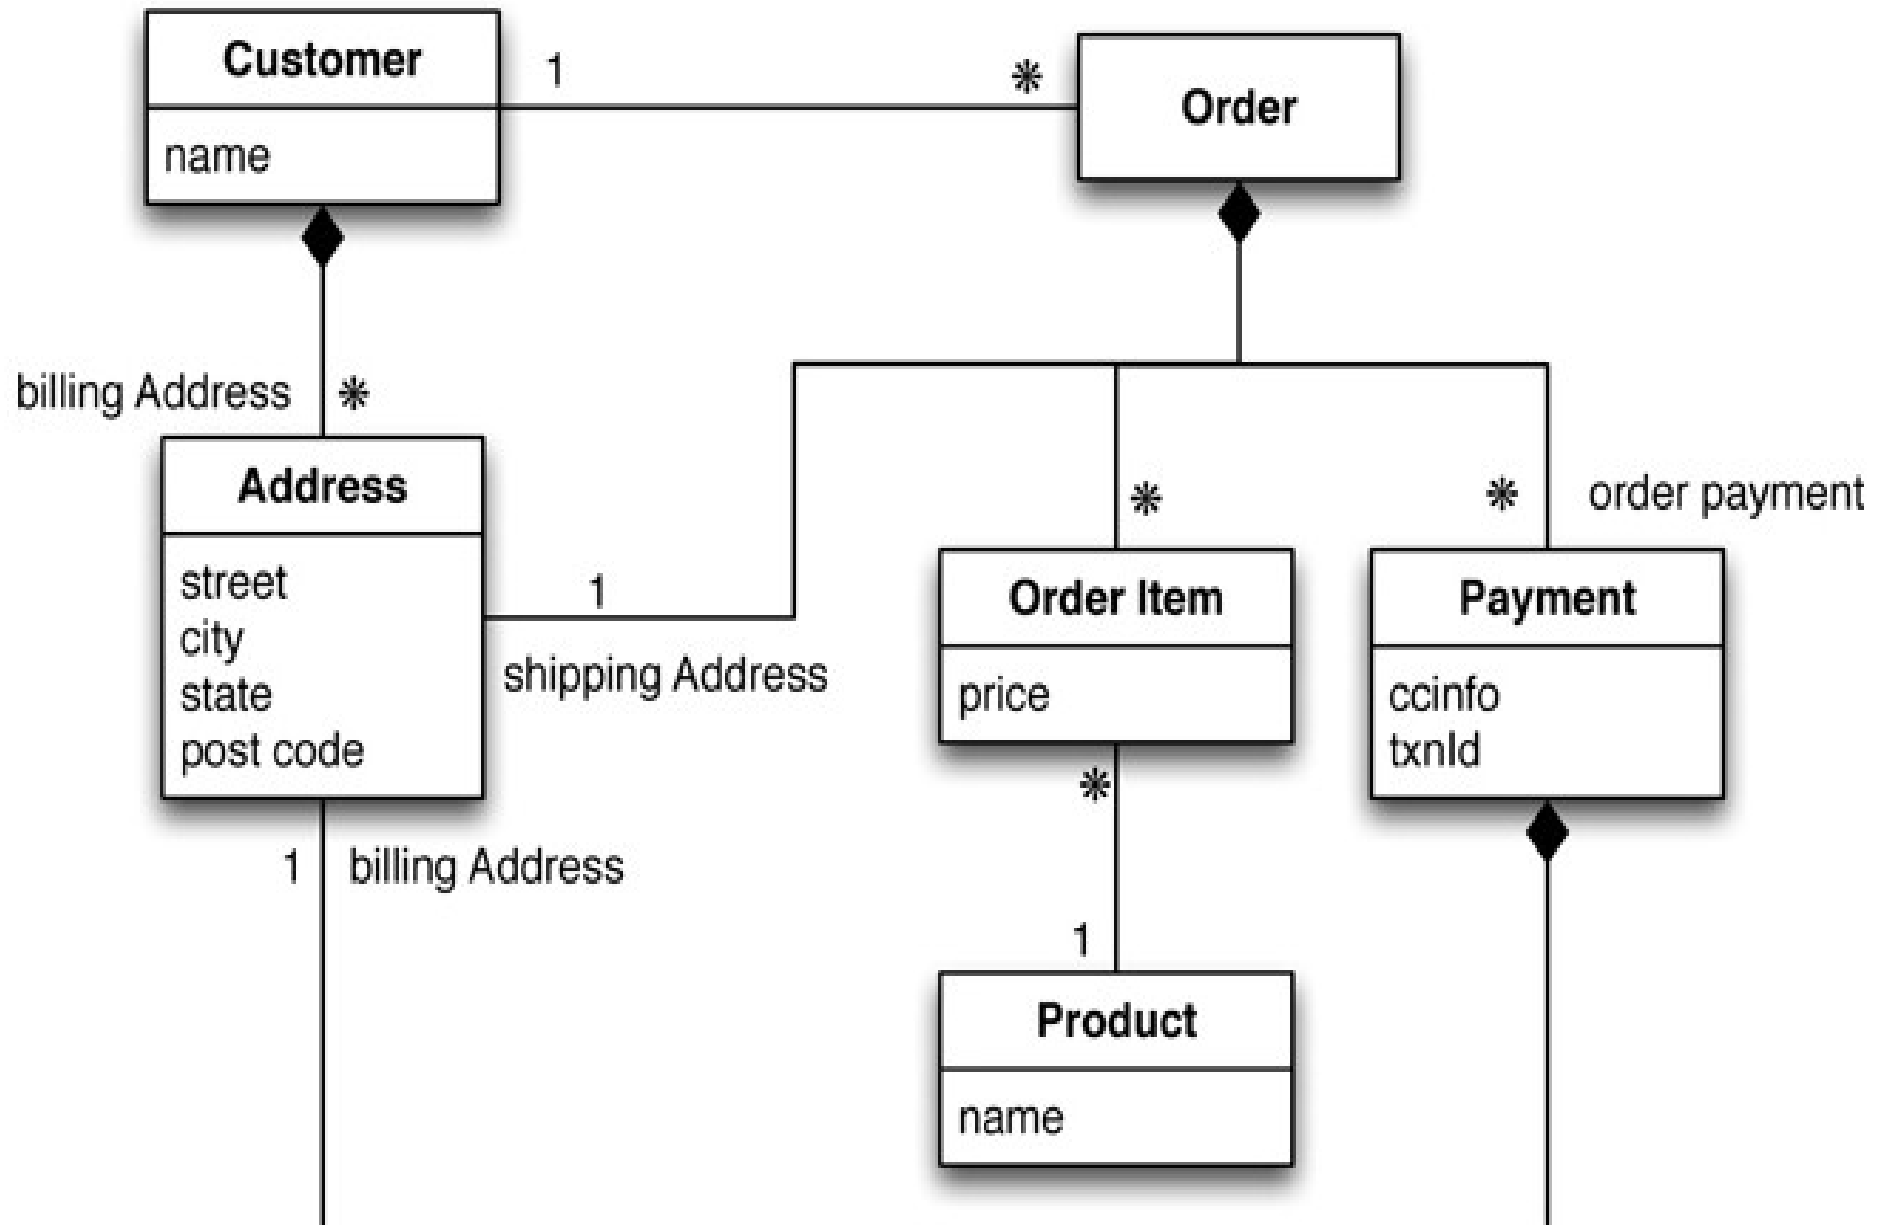
BillingAddress		
Id	CustomerId	AddressId
55	1	77

OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45

Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abelif879rft

An aggregate data model

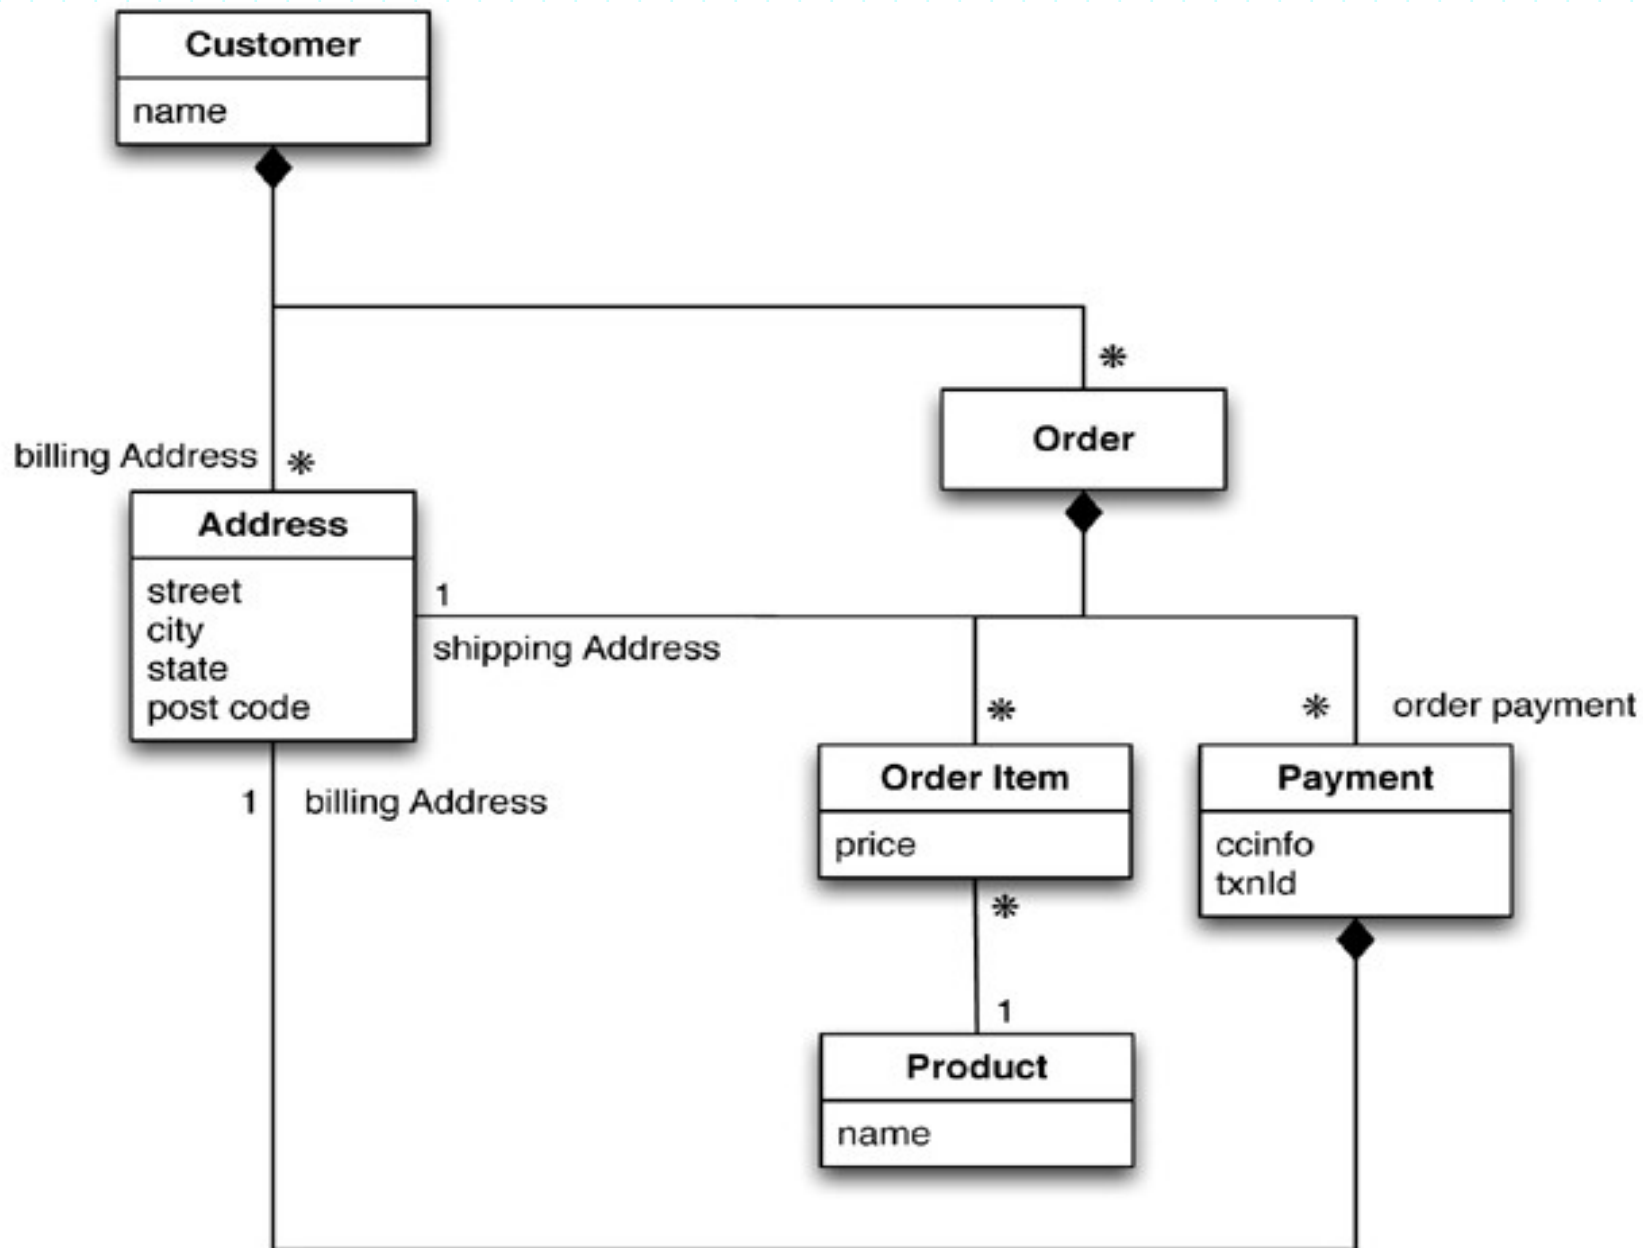


Aggregate Data

```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment":[
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnId":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}
```

A different aggregate data model



A different set of data

```
// in customers
{
  "customer": {
    "id": 1,
    "name": "Martin",
    "billingAddress": [{"city": "Chicago"}],
    "orders": [
      {
        "id": 99,
        "customerId": 1,
        "orderItems": [
          {
            "productId": 27,
            "price": 32.45,
            "productName": "NoSQL Distilled"
          }
        ]
      },
      {
        "id": 100,
        "customerId": 1,
        "orderItems": [
          {
            "productId": 28,
            "price": 19.99,
            "productName": "NoSQL Distilled"
          }
        ]
      }
    ],
    "shippingAddress": [{"city": "Chicago"}]
```

Consequences

- How are these aggregates represented in relational databases?
- What benefits are gained from using these aggregates in this way?
- ACID Transactions

Schemaless Databases

- Under this model we do not require a schema
 - How is this possible?
 - Advantages?
 - Disadvantages?

- Are they actually schemaless?

Modeling for Data Access

```
# Customer object
{
  "customerId": 1,
  "name": "Martin",
  "billingAddress": [{"city": "Chicago"}],
  "payment": [
    {"type": "debit",
     "ccinfo": "1000-1000-1000-1000"}
  ]
}

# Order object
{
  "orderId": 99,
  "customerId": 1,
  "orderDate": "Nov-20-2011",
  "orderItems": [{"productId": 27, "price": 32.45}],
  "orderPayment": [{"ccinfo": "1000-1000-1000-1000",
                    "txnId": "abelif879rft"}],
  "shippingAddress": {"city": "Chicago"}
}
```



Distribution

- One of the major benefits of NoSQL (why?)
 - Comes with one major drawback...

Single Server

- Is this really distribution?
- Why might you want to do it this way?

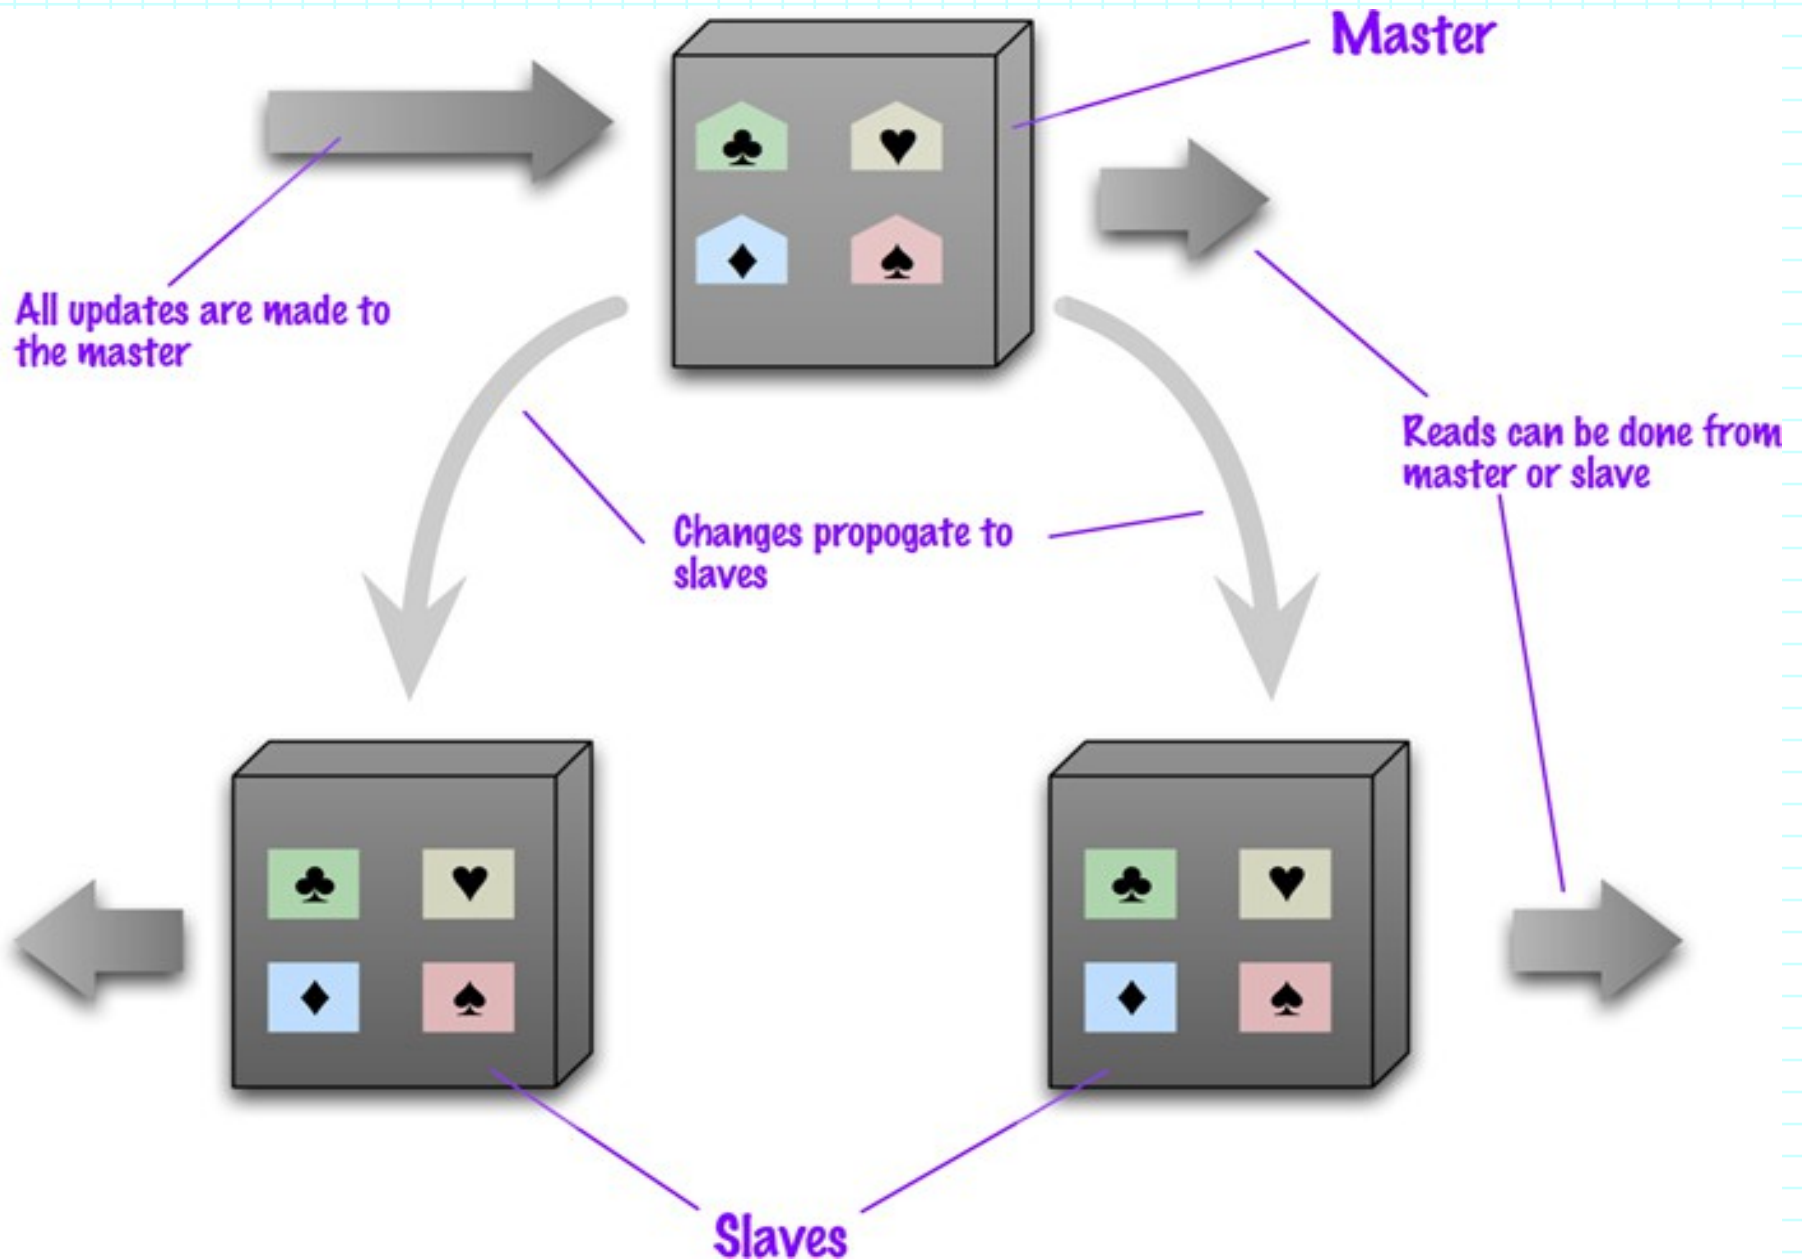
Sharding



Sharding

- What is the ideal case for # of users / servers?
- How do we decide how to split up the data?
 - What data will commonly be accessed together?
- What can we use to help us perform these tasks?


Master-Slave Replication



Master-Slave Replication

- How does this affect performance of reads? Writes?
- How does this affect data resilience?
- What about consistency?

Peer-to-Peer Replication




Peer-to-Peer Replication


- How does this affect performance of reads? Writes?
- How does this affect data resilience?
- What about consistency?

Combinations


master for two shards




slave for two shards



master for one shard



master for one shard
and slave for a shard



slave for two shards

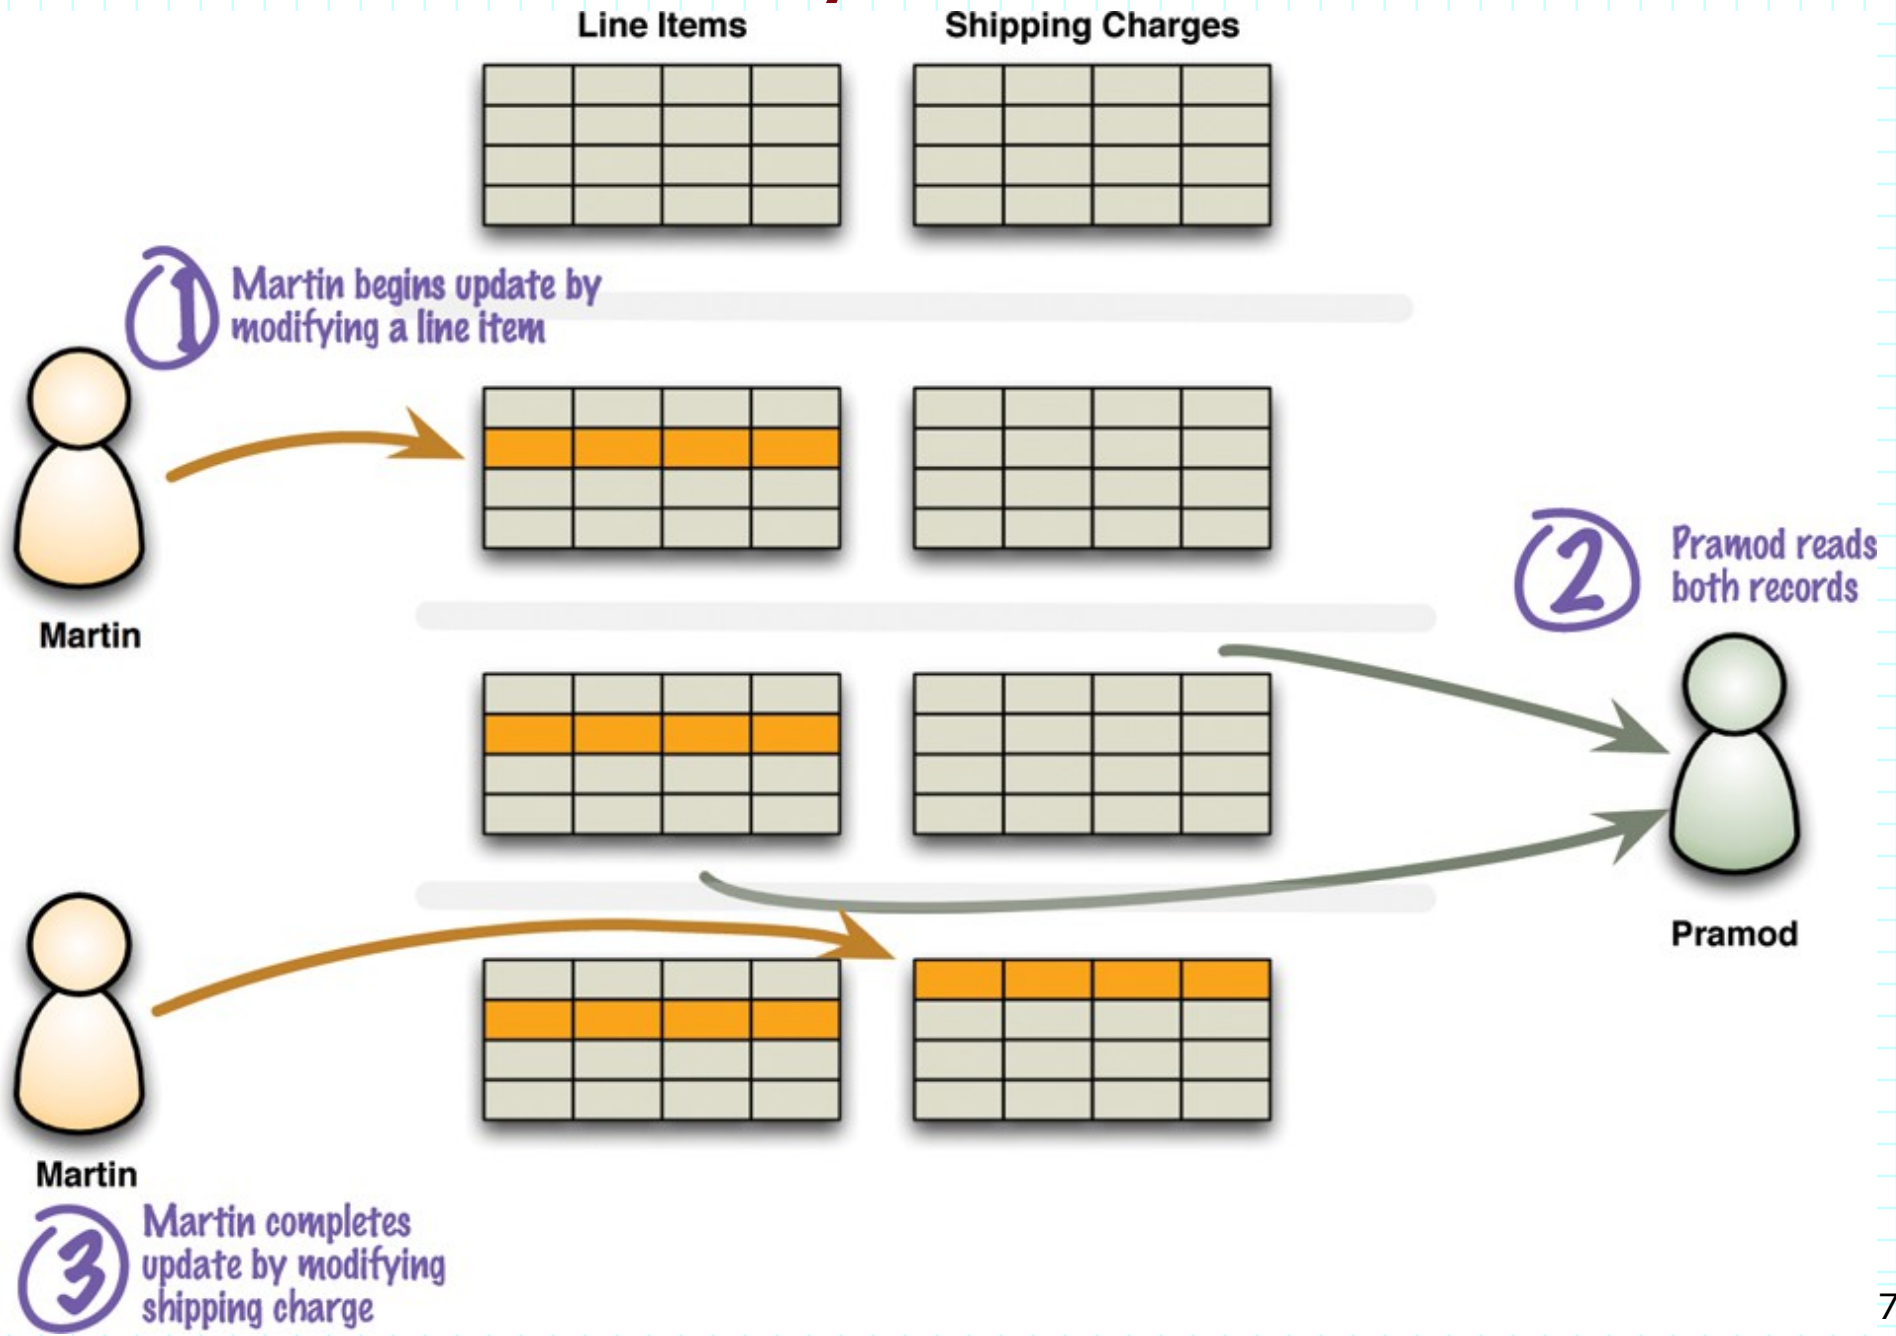


slave for one shard

Update Consistency

- Imagine two people try to update the same piece of data at the same time
 - What happens?
 - What do we want to happen?
- Pessimistic vs. Optimistic

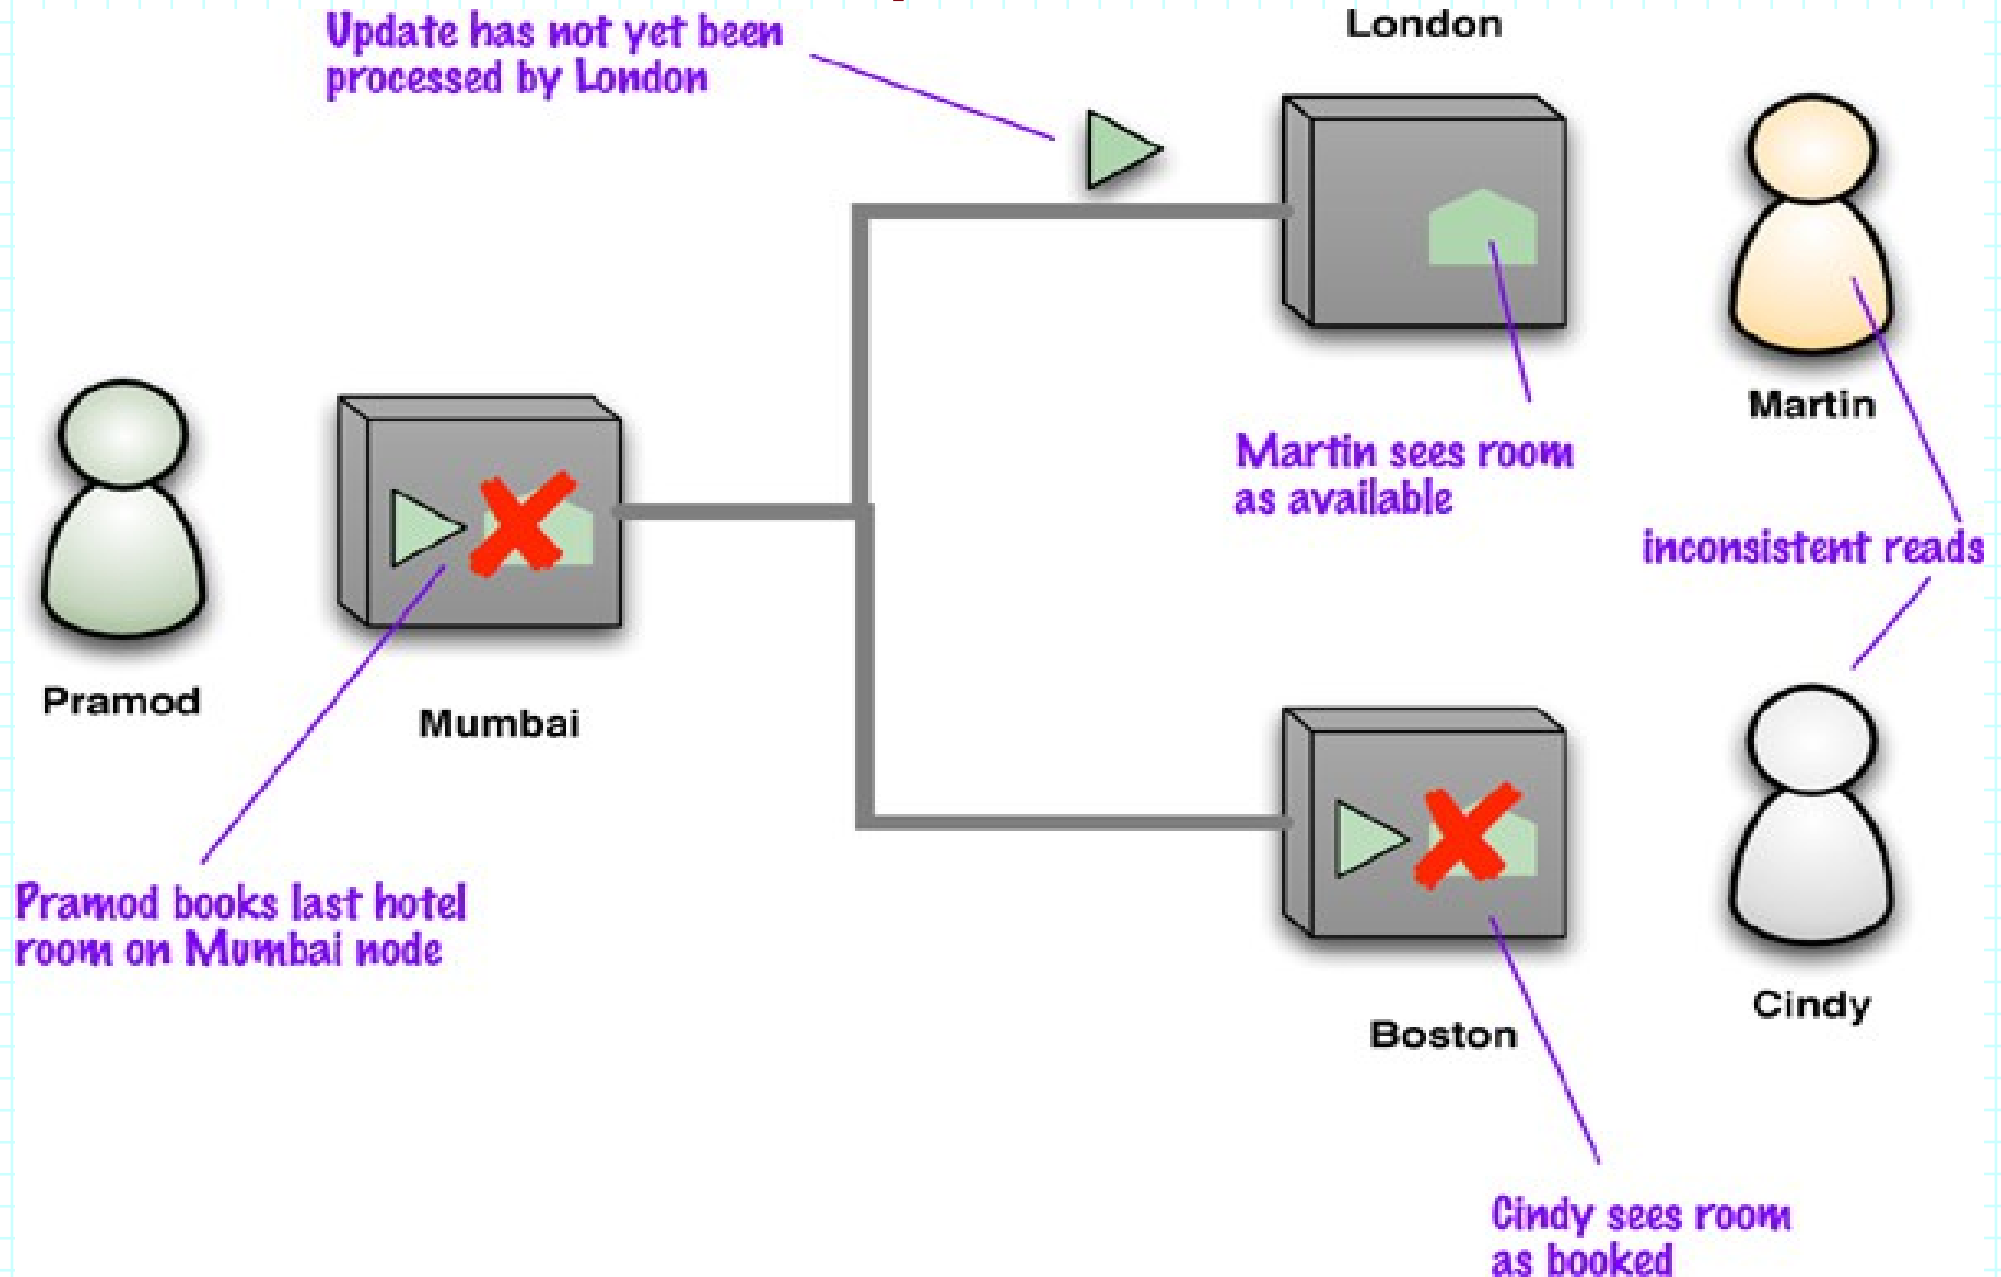
Read Consistency



Read Consistency

- Logical consistency
 - How do relational DBs handle this?
- How does NoSQL handle it?
 - Inconsistency window
- How does replication complicate this kind of consistency?

Read Consistency



Read Consistency

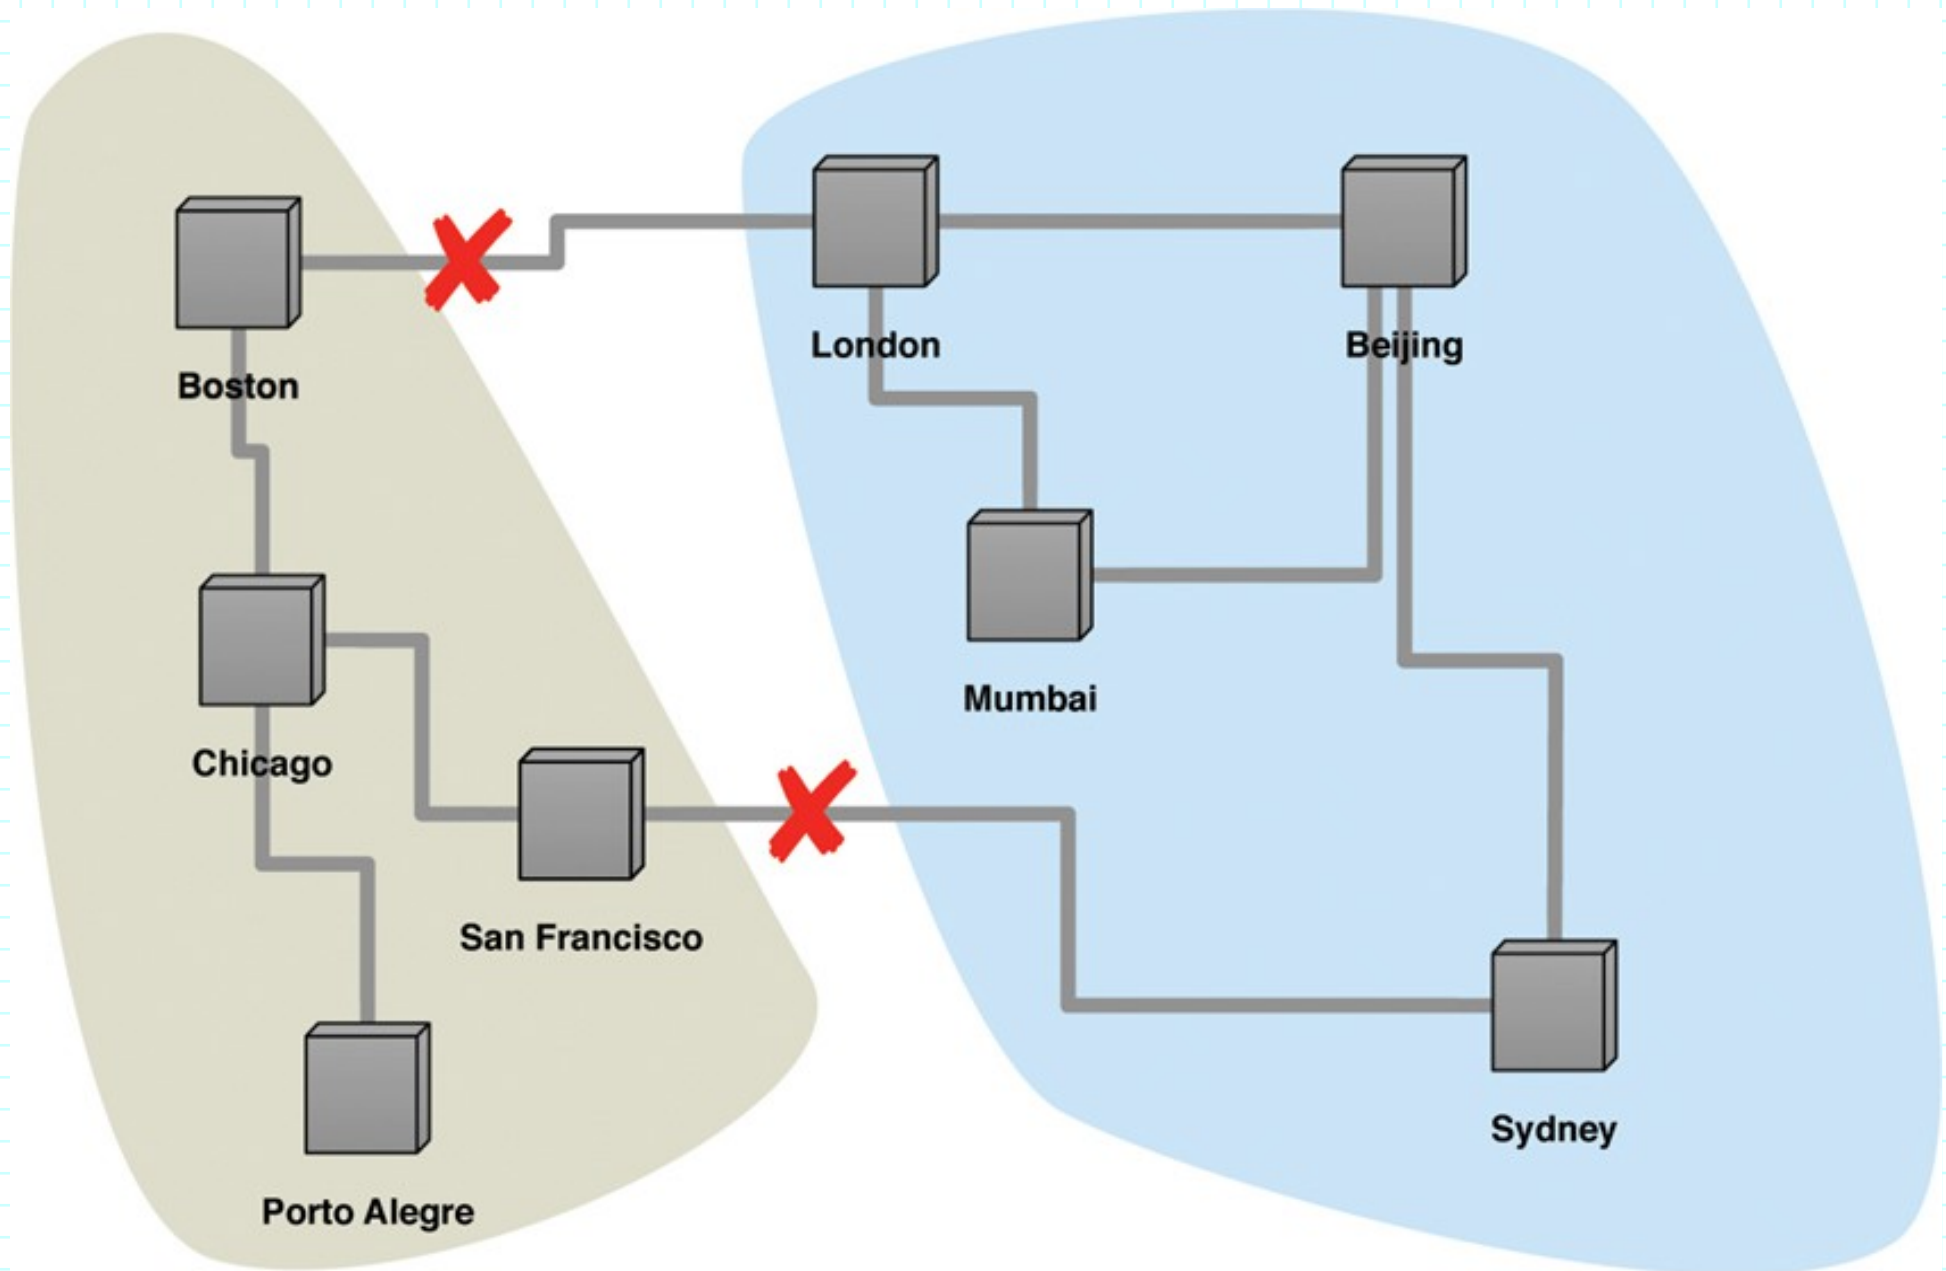
- Read-your-writes consistency
- Session consistency
 - Sticky Sessions

Relaxing Consistency

- It is always possible to design a system that is consistent
 - Why might we want to sacrifice this?

- CAP Theorem

Partition Tolerance



CAP Theorem

- What is an example of a system without partition tolerance?
 - Should we aim for this?
- It is all about the trade off!

CAP Theorem

- How can we improve consistency?
- How can we improve availability?
- Inconsistent writes
 - What about reads?

Quorums

- How many nodes do we actually need to get consistency?
- Write quorums
- Replication factor
- How does this affect reads?