

Database Management Systems

- Relational Operations

Retrieving information

- SELECT
- FROM
- WHERE
- ORDER BY
- GROUP BY

Relational Algebra

- All of these queries can be broken down using relational algebra
 - Operates on sets
- Unary
 - Select
 - Project
 - Rename
- Relational
 - Union
 - Intersection
 - Difference
 - Cartesian Product

Select

- Careful!
 - Not like SELECT
- Specifies which tuples to keep from a relation based on a condition:

$\sigma_{\text{hire date after 1-1-1990}}(\text{EMPLOYEES})$

Select

- Result set will have the same attributes as the relation
- Its commutative!
- Bound on number of tuples?

Project

- Keeps certain columns, discards the rest

Π LNAME, FNAME, SALARY (EMPLOYEE)

Roles of a DBMS

- Duplicate tuples are removed (why?)
- Bound on tuples?
- Commutative?

Rename

- Applies a different name to attribute(s) of a relation
 - Necessary in some cases

$\rho_{(A1, A2, A3, A4, A5)}(R)$

Union

$R \cup S$ includes all tuples in R , in S , or in both R and S

Relations must contain the same number of columns

Columns must have the same data type

SQL Syntax:

```
SELECT *  
FROM dept_emp  
UNION  
SELECT *  
FROM dept_manager
```

Intersection and Difference

$R \cap S$ includes all tuples that are in R and S
* INTERSECT

$R - S$ includes all tuples that are in R but not in S
Not often used

■ Commutative?

Cartesian Product

- Combine tuples from two relations combinatorially
 - $R \times S$
- Number of attributes in the result?
- Number of tuples in the result?

Cartesian Product

- Under what conditions is this a useful operation?
 - What do we have at our disposal to help us with this operation?

JOIN

- Cartesian Product combined with a select
- Example:

```
SELECT *  
FROM employees JOIN salaries ON  
    employees.emp_no = salaries.emp_no
```

- Can join more than 2 tables if necessary

Outer JOINS

- Three types
 - LEFT
 - RIGHT
 - FULL
- What are “missing” values replaced with?
- What could this be used for?

Combinations

- The six operations (select, project, union, difference, rename, cartesian product) are a complete set
 - Any other expression can be expressed using these operations
- Question: how can we express intersection using these operations?
- What order are the SQL clauses that we've been using executed in?

Aggregates

- “outside” the realm of relational algebra

$\mathcal{F}_{\text{AVERAGE SALARY}}$ (SALARIES)

$\text{emp_no } \mathcal{F}_{\text{AVERAGE SALARY}}$ (SALARIES)

Query Trees

- Data structure used to organize the operations to be performed
- Example: Employee names and salaries

Relational Calculus

- Relational algebra is procedural
- Relational calculus is declarative
 - No order of operations
- Tuple Calculus:
 $\{t.\text{first_name} \mid \text{EMPLOYEE}(t) \text{ AND } t.\text{birth_date} > 1960-01-01\}$
- Domain Calculus:
 $\{ uv \mid (\exists r) (\exists s) (\exists t)(\text{EMPLOYEE}(rstuv) \text{ and } r > 1960-01-01) \}$

Quantifiers

- \exists is called the existential qualifier. To satisfy the condition, a tuple must exist within the specified domain of tuples.
- \forall is called the universal qualifier - every tuple must conform to the condition.

Practice Problems

- Write a query to show the names and birthdays of all employees
 - What relational operations did you use?
 - Write the query using relational algebra
 - Write the query using relational (tuple) calculus

Practice Problems

- Write a query to show who the managers are of each employee
 - Hint: you'll need a join (maybe more than 2 tables?)
 - What relational operations did you use?
 - Write the query using relational algebra
 - Write the query using relational (tuple) calculus